

Time, Clocks, and the Ordering of Events in a Distributed System

Leslie Lamport (Massachusetts Computer Associates, Inc.)

The concept of one event happening before another in a distributed system is examined, and is shown to define a partial ordering of the events. A distributed algorithm is given for synchronizing a system of logical clocks which can be used to totally order the events. The use of the total ordering is illustrated with a method for solving synchronization problems. The algorithm is then specialized for synchronizing physical clocks, and a bound is derived on how far out of synchrony the clocks can become.

Key Words and Phrases: distributed systems, computer networks, clock synchronization, multiprocess systems

CR Categories: 4.32, 5.29

General permission to make fair use in teaching or research of all or part of this material is granted to individual readers and to nonprofit libraries acting for them provided that ACM's copyright notice is given and that reference is made to the publication, to its date of issue, and to the fact that reprinting privileges were granted by permission of the Association for Computing Machinery. To otherwise reprint a figure, table, other substantial excerpt, or the entire work requires specific permission as does republication, or systematic or multiple reproduction.

This work was supported by the Advanced Research Projects Agency of the Department of Defense and Rome Air Development Center. It was monitored by Rome Air Development Center under contract number F 30602-76-C-0094.

Author's address: Computer Science Laboratory, SRI International, 333 Ravenswood Ave., Menlo Park CA 94025.

© 1978 ACM 0001-0782/78/0700-0558 \$00.75

Paper originally published in *Communications of the ACM* 21(7), July 1978.

Introduction

The concept of time is fundamental to our way of thinking. It is derived from the more basic concept of the order in which events occur. We say that something happened at 3:15 if it occurred *after* our clock read 3:15 and before it read 3:16. The concept of the temporal ordering of events pervades our thinking about systems. For example, in an airline reservation system we specify that a request for a reservation should be granted if it is made *before* the flight is filled. However, we will see that this concept must be carefully reexamined when considering events in a distributed system.

A distributed system consists of a collection of distinct processes which are spatially separated, and which communicate with one another by exchanging messages. A network of interconnected computers, such as the ARPA net, is a distributed system. A single computer can also be viewed as a distributed system in which the central control unit, the memory units, and the input-output channels are separate processes. A system is distributed if the message transmission delay is not negligible compared to the time between events in a single process.

We will concern ourselves primarily with systems of spatially separated computers. However, many of our remarks will apply more generally. In particular, a multiprocessing system on a single computer involves problems similar to those of a distributed system because of the unpredictable order in which certain events can occur.

In a distributed system, it is sometimes impossible to say that one of two events occurred first. The relation “happened before” is therefore only a partial ordering of the events in the system. We have found that problems often arise because people are not fully aware of this fact and its implications.

In this paper, we discuss the partial ordering defined by the “happened before” relation, and give a distributed algorithm for extending it to a consistent total ordering of all the events. This algorithm can provide a useful mechanism for implementing a distributed system. We illustrate its use with a simple method for solving synchronization problems. Unexpected, anomalous behavior can occur if the ordering obtained by this algorithm differs from that perceived by the user. This can be avoided by introducing real, physical clocks. We describe a simple method for synchronizing these clocks, and derive an upper bound on how far out of synchrony they can drift.

The Partial Ordering

Most people would probably say that an event a happened before an event b if a happened at an earlier time than b . They might justify this definition in terms of

physical theories of time. However, if a system is to meet a specification correctly, then that specification must be given in terms of events observable within the system. If the specification is in terms of physical time, then the system must contain real clocks. Even if it does contain real clocks, there is still the problem that such clocks are not perfectly accurate and do not keep precise physical time. We will therefore define the “happened before” relation without using physical clocks.

We begin by defining our system more precisely. We assume that the system is composed of a collection of processes. Each process consists of a sequence of events. Depending upon the application, the execution of a subprogram on a computer could be one event, or the execution of a single machine instruction could be one event. We are assuming that the events of a process form a sequence, where a occurs before b in this sequence if a happens before b . In other words, a single process is defined to be a set of events with an a priori total ordering. This seems to be what is generally meant by a process.¹ It would be trivial to extend our definition to allow a process to split into distinct subprocesses, but we will not bother to do so.

We assume that sending or receiving a message is an event in a process. We can then define the “happened before” relation, denoted by “ \rightarrow ”, as follows.

Definition The relation “ \rightarrow ” on the set of events of a system is the smallest relation satisfying the following three conditions: (1) If a and b are events in the same process, and a comes before b , then $a \rightarrow b$. (2) If a is the sending of a message by one process and b is the receipt of the same message by another process, then $a \rightarrow b$. (3) If $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$. Two distinct events a and b are said to be *concurrent* if $a \not\rightarrow b$ and $b \not\rightarrow a$.

We assume that $a \not\rightarrow a$ for any event a . (Systems in which an event can happen before itself do not seem to be physically meaningful.) This implies that \rightarrow is an irreflexive partial ordering on the set of all events in the system.

It is helpful to view this definition in terms of a “space-time diagram” such as Figure 1. The horizontal direction represents space, and the vertical direction represents time—later times being higher than earlier ones. The dots denote events, the vertical lines denote processes, and the wavy lines denote messages.² It is easy

1. The choice of what constitutes an event affects the ordering of events in a process. For example, the receipt of a message might denote the setting of an interrupt bit in a computer, or the execution of a subprogram to handle that interrupt. Since interrupts need not be handled in the order that they occur, this choice will affect the ordering of a process’ message-receiving events.

2. Observe that messages may be received out of order. We allow the sending of several messages to be a single event, but for convenience we will assume that the receipt of a single message does not coincide with the sending or receipt of any other message.




Figure 1

to see that $a \rightarrow b$ means that one can go from a to b in the diagram by moving forward in time along process and message lines. For example, we have $p_1 \rightarrow r_4$ in Figure 1.

Another way of viewing the definition is to say that $a \rightarrow b$ means that it is possible for event a to causally affect event b . Two events are concurrent if neither can causally affect the other. For example, events p_3 and q_3 of Figure 1 are concurrent. Even though we have drawn the diagram to imply that q_3 occurs at an earlier physical time than p_3 , process P cannot know what process Q did at q_3 until it receives the message at p_4 . (Before event p_4 , P could at most know what Q was *planning* to do at q_3 .)

This definition will appear quite natural to the reader familiar with the invariant space-time formulation of special relativity, as described for example in [1] or the first chapter of [2]. In relativity, the ordering of events is defined in terms of messages that *could* be sent. However, we have taken the more pragmatic approach of only considering messages that actually *are* sent. We should be able to determine if a system performed correctly by knowing only those events which *did* occur, without knowing which events *could* have occurred.

Logical Clocks

We now introduce clocks into the system. We begin with an abstract point of view in which a clock is just a way of assigning a number to an event, where the number is thought of as the time at which the event occurred. More precisely, we define a

clock C_i for each process P_i to be a function which assigns a number $C_i(a)$ to any event a in that process. The entire system of clocks is represented by the function C which assigns to any event b the number $C(b)$, where $C(b) = C_j(b)$ if b is an event in process P_j . For now, we make no assumption about the relation of the numbers $C_i(a)$ to physical time, so we can think of the clocks C_i as logical rather than physical clocks. They may be implemented by counters with no actual timing mechanism.

We now consider what it means for such a system of clocks to be correct. We cannot base our definition of correctness on physical time, since that would require introducing clocks which keep physical time. Our definition must be based on the order in which events occur. The strongest reasonable condition is that if an event a occurs before another event b , then a should happen at an earlier time than b . We state this condition more formally as follows.

Clock Condition For any events a, b : if $a \rightarrow b$ then $C(a) < C(b)$.

Note that we cannot expect the converse condition to hold as well, since that would imply that any two concurrent events must occur at the same time. In Figure 1, p_2 and p_3 are both concurrent with q_3 , so this would mean that they both must occur at the same time as q_3 , which would contradict the Clock Condition because $p_2 \rightarrow p_3$.

It is easy to see from our definition of the relation “ \rightarrow ” that the Clock Condition is satisfied if the following two conditions hold.

- C1** If a and b are events in process P_i , and a comes before b , then $C_i(a) < C_i(b)$.
- C2** If a is the sending of a message by process P_i and b is the receipt of that message by process P_j , then $C_i(a) < C_j(b)$.

Let us consider the clocks in terms of a space-time diagram. We imagine that a process’ clock “ticks” through every number, with the ticks occurring between the process’ events. For example, if a and b are consecutive events in process P_i with $C_i(a) = 4$ and $C_i(b) = 7$, then clock ticks 5, 6, and 7 occur between the two events. We draw a dashed “tick line” through all the like-numbered ticks of the different processes. The space-time diagram of Figure 1 might then yield the picture in Figure 2. Condition C1 means that there must be a tick line between any two events on a process line, and condition C2 means that every message line must cross a tick line. From the pictorial meaning of \rightarrow , it is easy to see why these two conditions imply the Clock Condition.

We can consider the tick lines to be the time coordinate lines of some Cartesian coordinate system on space-time. We can redraw Figure 2 to straighten these




Figure 2




Figure 3

coordinate lines, thus obtaining Figure 3. Figure 3 is a valid alternate way of representing the same system of events as Figure 2. Without introducing the concept of physical time into the system (which requires introducing physical clocks), there is no way to decide which of these pictures is a better representation.

The reader may find it helpful to visualize a two-dimensional spatial network of processes, which yields a three-dimensional space-time diagram. Processes and messages are still represented by lines, but tick lines become two-dimensional surfaces.

Let us now assume that the processes are algorithms, and the events represent certain actions during their execution. We will show how to introduce clocks into the processes which satisfy the Clock Condition. Process P_i 's clock is represented by a register C_i , so that $C_i(a)$ is the value contained by C_i during the event a . The value of C_i will change between events, so changing C_i does not itself constitute an event.

To guarantee that the system of clocks satisfies the Clock Condition, we will insure that it satisfies conditions C1 and C2. Condition C1 is simple; the processes need only obey the following implementation rule:

- IR1** Each process P_i increments C_i between any two successive events.

To meet condition C2, we require that each message m contain a *timestamp* T_m which equals the time at which the message was sent. Upon receiving a message timestamped T_m , a process must advance its clock to be later than T_m . More precisely, we have the following rule.

- IR2** (a) If event a is the sending of a message m by process P_i , then the message m contains a timestamp $T_m = C_i(a)$. (b) Upon receiving a message m , process P_j sets C_j greater than or equal to its present value and greater than T_m .

In IR2(b) we consider the event which represents the receipt of the message m to occur after the setting of C_j . (This is just a notational nuisance, and is irrelevant in any actual implementation.) Obviously, IR2 insures that C2 is satisfied. Hence, the simple implementation rules IR1 and IR2 imply that the Clock Condition is satisfied, so they guarantee a correct system of logical clocks.

Ordering the Events Totally

We can use a system of clocks satisfying the Clock Condition to place a total ordering on the set of all system events. We simply order the events by the times at which they occur. To break ties, we use any arbitrary total ordering $<$ of the processes. More precisely, we define a relation \Rightarrow as follows: if a is an event in process P_i and b is an event in process P_j , then $a \Rightarrow b$ if and only if either (i) $C_i(a) < C_j(b)$ or (ii) $C_i(a) = C_j(b)$ and $P_i < P_j$. It is easy to see that this defines a total ordering, and that the Clock Condition implies that if $a \rightarrow b$ then $a \Rightarrow b$. In

other words, the relation \Rightarrow is a way of completing the “happened before” partial ordering to a total ordering.³

The ordering \Rightarrow depends upon the system of clocks C_i , and is not unique. Different choices of clocks which satisfy the Clock Condition yield different relations \Rightarrow . Given any total ordering relation \Rightarrow which extends \rightarrow , there is a system of clocks satisfying the Clock Condition which yields that relation. It is only the partial ordering \rightarrow which is uniquely determined by the system of events.

Being able to totally order the events can be very useful in implementing a distributed system. In fact, the reason for implementing a correct system of logical clocks is to obtain such a total ordering. We will illustrate the use of this total ordering of events by solving the following version of the mutual exclusion problem. Consider a system composed of a fixed collection of processes which share a single resource. Only one process can use the resource at a time, so the processes must synchronize themselves to avoid conflict. We wish to find an algorithm for granting the resource to a process which satisfies the following three conditions: (I) A process which has been granted the resource must release it before it can be granted to another process. (II) Different requests for the resource must be granted in the order in which they are made. (III) If every process which is granted the resource eventually releases it, then every request is eventually granted.

We assume that the resource is initially granted to exactly one process.

These are perfectly natural requirements. They precisely specify what it means for a solution to be correct.⁴ Observe how the conditions involve the ordering of events. Condition II says nothing about which of two concurrently issued requests should be granted first.

It is important to realize that this is a nontrivial problem. Using a central scheduling process which grants requests in the order they are received will not work, unless additional assumptions are made. To see this, let P_0 be the scheduling process. Suppose P_1 sends a request to P_0 and then sends a message to P_2 . Upon receiving the latter message, P_2 sends a request to P_0 . It is possible for P_2 's request to reach P_0 before P_1 's request does. Condition II is then violated if P_2 's request is granted first.

To solve the problem, we implement a system of clocks with rules IR1 and IR2, and use them to define a total ordering \Rightarrow of all events. This provides a total

3. The ordering $<$ establishes a priority among the processes. If a “fairer” method is desired, then $<$ can be made a function of the clock value. For example, if $C_i(a) = C_j(b)$ and $j < i$, then we can let $a \Rightarrow b$ if $j < C_i(a) \bmod N \leq i$, and $b \Rightarrow a$ otherwise; where N is the total number of processes.

4. The term “eventually” should be made precise, but that would require too long a diversion from our main topic.

ordering of all request and release operations. With this ordering, finding a solution becomes a straightforward exercise. It just involves making sure that each process learns about all other processes' operations.

To simplify the problem, we make some assumptions. They are not essential, but they are introduced to avoid distracting implementation details. We assume first of all that for any two processes P_i and P_j , the messages sent from P_i to P_j are received in the same order as they are sent. Moreover, we assume that every message is eventually received. (These assumptions can be avoided by introducing message numbers and message acknowledgment protocols.) We also assume that a process can send messages directly to every other process.

Each process maintains its own *request queue* which is never seen by any other process. We assume that the request queues initially contain the single message $T_0: P_0 \text{ requests resource}$, where P_0 is the process initially granted the resource and T_0 is less than the initial value of any clock.

The algorithm is then defined by the following five rules. For convenience, the actions defined by each rule are assumed to form a single event.

1. To request the resource, process P_i sends the message $T_m: P_i \text{ requests resource}$ to every other process, and puts that message on its request queue, where T_m is the timestamp of the message.
2. When process P_j receives the message $T_m: P_i \text{ requests resource}$, it places it on its request queue and sends a (timestamped) acknowledgment message to P_i .⁵
3. To release the resource, process P_i removes any $T_m: P_i \text{ requests resource}$ message from its request queue and sends a (timestamped) $P_i \text{ releases resource}$ message to every other process.
4. When process P_j receives a $P_i \text{ releases resource}$ message, it removes any $T_m: P_i \text{ requests resource}$ message from its request queue.
5. Process P_i is granted the resource when the following two conditions are satisfied: (i) There is a $T_m: P_i \text{ requests resource}$ message in its request queue which is ordered before any other request in its queue by the relation \Rightarrow . (To define the relation " \Rightarrow " for messages, we identify a message with the event of sending it.) (ii) P_i has received a message from every other process timestamped later than T_m .⁶

5. This acknowledgment message need not be sent if P_j has already sent a message to P_i timestamped later than T_m .

6. If $P_i < P_j$, then P_i need only have received a message timestamped $\geq T_m$ from P_j .

Note that conditions (i) and (ii) of rule 5 are tested locally by P_i .

It is easy to verify that the algorithm defined by these rules satisfies conditions I–III. First of all, observe that condition (ii) of rule 5, together with the assumption that messages are received in order, guarantees that P_i has learned about all requests which preceded its current request. Since rules 3 and 4 are the only ones which delete messages from the request queue, it is then easy to see that condition I holds. Condition II follows from the fact that the total ordering \Rightarrow extends the partial ordering \rightarrow . Rule 2 guarantees that after P_i requests the resource, condition (ii) of rule 5 will eventually hold. Rules 3 and 4 imply that if each process which is granted the resource eventually releases it, then condition (i) of rule 5 will eventually hold, thus proving condition III.

This is a distributed algorithm. Each process independently follows these rules, and there is no central synchronizing process or central storage. This approach can be generalized to implement any desired synchronization for such a distributed multiprocess system. The synchronization is specified in terms of a *State Machine*, consisting of a set C of possible commands, a set S of possible states, and a function $e: C \times S \rightarrow S$. The relation $e(C, S) = S'$ means that executing the command C with the machine in state S causes the machine state to change to S' . In our example, the set C consists of all the commands P_i *requests resource* and P_i *releases resource*, and the state consists of a queue of waiting request commands, where the request at the head of the queue is the currently granted one. Executing a *request* command adds the request to the tail of the queue, and executing a *release* command removes a command from the queue.⁷

Each process independently simulates the execution of the State Machine, using the commands issued by all the processes. Synchronization is achieved because all processes order the commands according to their timestamps (using the relation \Rightarrow), so each process uses the same sequence of commands. A process can execute a command timestamped T when it has learned of all commands issued by all other processes with timestamps less than or equal to T . The precise algorithm is straightforward, and we will not bother to describe it.

This method allows one to implement any desired form of multiprocess synchronization in a distributed system. However, the resulting algorithm requires the active participation of all the processes. A process must know all the commands issued by other processes, so that the failure of a single process will make it impos-

7. If each process does not strictly alternate request and release commands, then executing a *release* command could delete zero, one, or more than one request from the queue.

sible for any other process to execute State Machine commands, thereby halting the system.

The problem of failure is a difficult one, and it is beyond the scope of this paper to discuss it in any detail. We will just observe that the entire concept of failure is only meaningful in the context of physical time. Without physical time, there is no way to distinguish a failed process from one which is just pausing between events. A user can tell that a system has “crashed” only because he has been waiting too long for a response. A method which works despite the failure of individual processes or communication lines is described in [3].

Anomalous Behavior

Our resource scheduling algorithm ordered the requests according to the total ordering \Rightarrow . This permits the following type of “anomalous behavior.” Consider a nationwide system of interconnected computers. Suppose a person issues a request A on a computer A, and then telephones a friend in another city to have him issue a request B on a different computer B. It is quite possible for request B to receive a lower timestamp and be ordered before request A. This can happen because the system has no way of knowing that A actually preceded B, since that precedence information is based on messages external to the system.

Let us examine the source of the problem more closely. Let S be the set of all system events. Let us introduce a set of events which contains the events in S together with all other relevant external events, such as the phone calls in our example. Let \rightarrow denote the “happened before” relation for \underline{S} . In our example, we had $A \rightarrow B$, but $A \not\rightarrow B$. It is obvious that no algorithm based entirely upon events in S , and which does not relate those events in any way with the other events in \underline{S} , can guarantee that request A is ordered before request B.

There are two possible ways to avoid such anomalous behavior. The first way is to explicitly introduce into the system the necessary information about the ordering \rightarrow . In our example, the person issuing request A could receive the timestamp T_A of that request from the system. When issuing request B, his friend could specify that B be given a timestamp later than T_A . This gives the user the responsibility for avoiding anomalous behavior.

The second approach is to construct a system of clocks which satisfies the following condition.

Strong Clock Condition For any events a, b in S : if $a \rightarrow b$ then $C(a) < C(b)$.

This is stronger than the ordinary Clock Condition because \rightarrow is a stronger relation than \rightarrow . It is not in general satisfied by our logical clocks.

Let us identify \underline{S} with some set of “real” events in physical space-time, and let \rightarrow be the partial ordering of events defined by special relativity. One of the mysteries of the universe is that it is possible to construct a system of physical clocks which, running quite independently of one another, will satisfy the Strong Clock Condition. We can therefore use physical clocks to eliminate anomalous behavior. We now turn our attention to such clocks.

Physical Clocks

Let us introduce a physical time coordinate into our space-time picture, and let $C_i(t)$ denote the reading of the clock C_i at physical time t .⁸ For mathematical convenience, we assume that the clocks run continuously rather than in discrete “ticks.” (A discrete clock can be thought of as a continuous one in which there is an error of up to $1/2$ “tick” in reading it.) More precisely, we assume that $C_i(t)$ is a continuous, differentiable function of t except for isolated jump discontinuities where the clock is reset. Then $dC_i(t)/dt$ represents the rate at which the clock is running at time t .

In order for the clock C_i to be a true physical clock, it must run at approximately the correct rate. That is, we must have $dC_i(t)/dt \approx 1$ for all t . More precisely, we will assume that the following condition is satisfied:

PC1 There exists a constant $\kappa \ll 1$ such that for all i : $|dC_i(t)/dt - 1| < \kappa$.

For typical crystal controlled clocks, $\kappa \leq 10^{-6}$.

It is not enough for the clocks individually to run at approximately the correct rate. They must be synchronized so that $C_i(t) \approx C_j(t)$ for all i, j , and t . More precisely, there must be a sufficiently small constant ϵ so that the following condition holds:

PC2 For all i, j : $|C_i(t) - C_j(t)| < \epsilon$.

If we consider vertical distance in Figure 2 to represent physical time, then PC2 states that the variation in height of a single tick line is less than ϵ .

Since two different clocks will never run at exactly the same rate, they will tend to drift further and further apart. We must therefore devise an algorithm to insure

8. We will assume a Newtonian space-time. If the relative motion of the clocks or gravitational effects are not negligible, then $C_i(t)$ must be deduced from the actual clock reading by transforming from proper time to the arbitrarily chosen time coordinate.

that PC2 always holds. First, however, let us examine how small κ and ϵ must be to prevent anomalous behavior. We must insure that the system $\underline{\mathcal{S}}$ of relevant physical events satisfies the Strong Clock Condition. We assume that our clocks satisfy the ordinary Clock Condition, so we need only require that the Strong Clock Condition holds when a and b are events in $\underline{\mathcal{S}}$ with $a \not\rightarrow b$. Hence, we need only consider events occurring in different processes.

Let μ be a number such that if event a occurs at physical time t and event b in another process satisfies $a \rightarrow b$, then b occurs later than physical time $t + \mu$. In other words, μ is less than the shortest transmission time for interprocess messages. We can always choose μ equal to the shortest distance between processes divided by the speed of light. However, depending upon how messages in $\underline{\mathcal{S}}$ are transmitted, μ could be significantly larger.

To avoid anomalous behavior, we must make sure that for any i, j , and t : $C_i(t + \mu) - C_j(t) > 0$. Combining this with PC1 and 2 allows us to relate the required smallness of κ and ϵ to the value of μ as follows. We assume that when a clock is reset, it is always set forward and never back. (Setting it back could cause C1 to be violated.) PC1 then implies that $C_i(t + \mu) - C_i(t) > (1 - \kappa)\mu$. Using PC2, it is then easy to deduce that $C_i(t + \mu) - C_j(t) > 0$ if the following inequality holds:

$$\epsilon/(1 - \kappa) \leq \mu.$$

This inequality together with PC1 and PC2 implies that anomalous behavior is impossible.

We now describe our algorithm for insuring that PC2 holds. Let m be a message which is sent at physical time t and received at time t' . We define $v_m = t' - t$ to be the *total delay* of the message m . This delay will, of course, not be known to the process which receives m . However, we assume that the receiving process knows some *minimum delay* $\mu_m \geq 0$ such that $\mu_m \leq v_m$. We call $\xi_m = v_m - \mu_m$ the *unpredictable delay* of the message.

We now specialize rules IR1 and 2 for our physical clocks as follows:

- IR1'** For each i , if P_i does not receive a message at physical time t , then C_i is differentiable at t and $dC_i(t)/dt > 0$.
- IR2'** (a) If P_i sends a message m at physical time t , then m contains a timestamp $T_m = C_i(t)$. (b) Upon receiving a message m at time t' , process P_j sets $C_j(t')$ equal to $\text{maximum}(C_j(t' - 0), T_m + \mu_m)$.⁹

9. $C_j(t' - 0) = \lim_{\delta \rightarrow 0} C_j(t' - |\delta|)$.

Although the rules are formally specified in terms of the physical time parameter, a process only needs to know its own clock reading and the timestamps of messages it receives. For mathematical convenience, we are assuming that each event occurs at a precise instant of physical time, and different events in the same process occur at different times. These rules are then specializations of rules IR1 and IR2, so our system of clocks satisfies the Clock Condition. The fact that real events have a finite duration causes no difficulty in implementing the algorithm. The only real concern in the implementation is making sure that the discrete clock ticks are frequent enough so C1 is maintained.

We now show that this clock synchronizing algorithm can be used to satisfy condition PC2. We assume that the system of processes is described by a directed graph in which an arc from process P_i to process P_j represents a communication line over which messages are sent directly from P_i to P_j . We say that a message is sent over this arc every τ seconds if for any t , P_i sends at least one message to P_j between physical times t and $t + \tau$. The *diameter* of the directed graph is the smallest number d such that for any pair of distinct processes P_j, P_k , there is a path from P_j to P_k having at most d arcs.

In addition to establishing PC2, the following theorem bounds the length of time it can take the clocks to become synchronized when the system is first started.

Theorem Assume a strongly connected graph of processes with diameter d which always obeys rules IR1' and IR2'. Assume that for any message m , $\mu_m \leq \mu$ for some constant μ , and that for all $t \geq t_0$: (a) PC1 holds. (b) There are constants τ and ξ such that every τ seconds a message with an unpredictable delay less than ξ is sent over every arc. Then PC2 is satisfied with $\epsilon \approx d(2\kappa\tau + \xi)$ for all $t \gtrsim t_0 + \tau d$, where the approximations assume $\mu + \xi \ll \tau$.

The proof of this theorem is surprisingly difficult, and is given in the Appendix. There has been a great deal of work done on the problem of synchronizing physical clocks. We refer the reader to [4] for an introduction to the subject. The methods described in the literature are useful for estimating the message delays μ_m and for adjusting the clock frequencies dC_i/dt (for clocks which permit such an adjustment). However, the requirement that clocks are never set backwards seems to distinguish our situation from ones previously studied, and we believe this theorem to be a new result.

Conclusion

We have seen that the concept of “happening before” defines an invariant partial ordering of the events in a distributed multiprocess system. We described an al-

gorithm for extending that partial ordering to a somewhat arbitrary total ordering, and showed how this total ordering can be used to solve a simple synchronization problem. A future paper will show how this approach can be extended to solve any synchronization problem.

The total ordering defined by the algorithm is somewhat arbitrary. It can produce anomalous behavior if it disagrees with the ordering perceived by the system's users. This can be prevented by the use of properly synchronized physical clocks. Our theorem showed how closely the clocks can be synchronized.

In a distributed system, it is important to realize that the order in which events occur is only a partial ordering. We believe that this idea is useful in understanding any multiprocess system. It should help one to understand the basic problems of multiprocessing independently of the mechanisms used to solve them.

Appendix

Proof of the Theorem

For any i and t , let us define C_i^t to be a clock which is set equal to C_i at time t and runs at the same rate as C_i , but is never reset. In other words,

$$C_i^t(t') = C_i(t) + \int_t^{t'} [dC_i(t)/dt]dt \quad (1)$$

for all $t' \geq t$. Note that

$$C_i(t') \geq C_i^t(t') \quad \text{for all } t' \geq t. \quad (2)$$

Suppose process P_1 at time t_1 sends a message to process P_2 which is received at time t_2 with an unpredictable delay $\leq \xi$, where $t_0 \leq t_1 \leq t_2$. Then for all $t \geq t_2$ we have:

$$\begin{aligned} C_2^{t_2}(t) &\geq C_2^{t_2}(t_2) + (1 - \kappa)(t - t_2) && [\text{by (1) and PC1}] \\ &\geq C_1(t_1) + \mu_m + (1 - \kappa)(t - t_2) && [\text{by IR2'(b)}] \\ &= C_1(t_1) + (1 - \kappa)(t - t_1) - [(t_2 - t_1) - \mu_m] + \kappa(t_2 - t_1) \\ &\geq C_1(t_1) + (1 - \kappa)(t - t_1) - \xi. \end{aligned}$$

Hence, with these assumptions, for all $t \geq t_2$ we have:

$$C_2^{t_2}(t) \geq C_1(t_1) + (1 - \kappa)(t - t_1) - \xi. \quad (3)$$

Now suppose that for $i = 1, \dots, n$ we have $t_i \leq t'_i < t_{i+1}$, $t_0 \leq t_1$, and that at time t'_i process P_i sends a message to process P_{i+1} which is received at time t_{i+1} with an unpredictable delay less than ξ . Then repeated application of the inequality (3)

yields the following result for $t \geq t_{n+1}$.

$$C_{n+1}^{t_{n+1}}(t) \geq C_1(t'_1) + (1 - \kappa)(t - t'_1) - n\xi. \quad (4)$$

From PC1, IRL' and 2' we deduce that

$$C_1(t'_1) \geq C_1(t_1) + (1 - \kappa)(t'_1 - t_1).$$

Combining this with (4) and using (2), we get

$$C_{n+1}(t) \geq C_1(t_1) + (1 - \kappa)(t - t_1) - n\xi \quad (5)$$

for $t \geq t_{n+1}$.

For any two processes P and P' , we can find a sequence of processes $P = P_0, P_1, \dots, P_{n+1} = P'$, $n \leq d$, with communication arcs from each P_i to P_{i+1} . By hypothesis (b) we can find times t_i, t'_i with $t'_i - t_i \leq \tau$ and $t_{i+1} - t'_i \leq \nu$, where $\nu = \mu + \xi$. Hence, an inequality of the form (5) holds with $n \leq d$ whenever $t \geq t_1 + d(\tau + \nu)$. For any i, j and any t, t_1 with $t_1 \geq t_0$ and $t \geq t_1 + d(\tau + \nu)$ we therefore have:

$$C_i(t) \geq C_j(t_1) + (1 - \kappa)(t - t_1) - d\xi. \quad (6)$$

Now let m be any message timestamped T_m , and suppose it is sent at time t and received at time t' . We pretend that m has a clock C_m which runs at a constant rate such that $C_m(t) = t_m$ and $C_m(t') = t_m + \mu_m$. Then $\mu_m \leq t' - t$ implies that $dC_m/dt \leq 1$. Rule IR2'(b) simply sets $C_j(t')$ to $\max(C_j(t' - 0), C_m(t'))$. Hence, clocks are reset only by setting them equal to other clocks.

For any time $t_x \geq t_0 + \mu/(1 - \kappa)$, let C_x be the clock having the largest value at time t_x . Since all clocks run at a rate less than $1 + \kappa$, we have for all i and all $t \geq t_x$:

$$C_i(t) \leq C_x(t_x) + (1 + \kappa)(t - t_x). \quad (7)$$

We now consider the following two cases: (i) C_x is the clock C_q of process P_q . (ii) C_x is the clock C_m of a message sent at time t_1 by process P_q . In case (i), (7) simply becomes

$$C_i(t) \leq C_q(t_x) + (1 + \kappa)(t - t_x). \quad (8i)$$

In case (ii), since $C_m(t_1) = C_q(t_1)$ and $dC_m/dt \leq 1$, we have

$$C_x(t_x) \leq C_q(t_1) + (t_x - t_1).$$

Hence, (7) yields

$$C_i(t) \leq C_q(t_1) + (1 + \kappa)(t - t_1). \quad (8ii)$$

Since $t_x \geq t_0 + \mu/(1 - \kappa)$, we get

$$\begin{aligned}
 C_q(t_x - \mu/(1 - \kappa)) &\leq C_q(t_x) - \mu && [\text{by PC1}] \\
 &\leq C_m(t_x) - \mu && [\text{by choice of } m] \\
 &\leq C_m(t_x) - (t_x - t_1)\mu_m/\nu_m && [\mu_m \leq \mu, t_x - t_1 \leq \nu_m] \\
 &= T_m && [\text{by definition of } C_m] \\
 &= C_q(t_1) && [\text{by IR2'(a)}]
 \end{aligned}$$

Hence, $C_q(t_x - \mu/(1 - \kappa)) \leq C_q(t_1)$, so $t_x - t_1 \leq \mu/(1 - \kappa)$ and thus $t_1 \geq t_0$.

Letting $t_1 = t_x$ in case (i), we can combine (8i) and (8ii) to deduce that for any t, t_x , with $t \geq t_x \geq t_0 + \mu/(1 - \kappa)$ there is a process P_q and a time t_1 with $t_x - \mu/(1 - \kappa) \leq t_1 \leq t_x$ such that for all i :

$$C_i(t) \leq C_q(t_1) + (1 + \kappa)(t - t_1). \quad (9)$$

Choosing t and t_x with $t \geq t_x + d(\tau + \nu)$, we can combine (6) and (9) to conclude that there exists a t_1 and a process P_q such that for all i :

$$\begin{aligned}
 C_q(t_1) + (1 - \kappa)(t - t_1) - d\xi &\leq C_i(t) \\
 &\leq C_q(t_1) + (1 + \kappa)(t - t_1)
 \end{aligned} \quad (10)$$

Letting $t = t_x + d(\tau + \nu)$, we get

$$d(\tau + \nu) \leq t - t_1 \leq d(\tau + \nu) + \mu/(1 - \kappa).$$

Combining this with (10), we get

$$\begin{aligned}
 C_q(t_1) + (t - t_1) - \kappa d(\tau + \nu) - d\xi &\leq C_i(t) \\
 &\leq C_q(t_1) + (t - t_1) + \kappa[d(\tau + \nu) + \mu/(1 - \kappa)]
 \end{aligned} \quad (11)$$

Using the hypothesis that $\kappa \ll 1$ and $\mu \leq \nu \ll \tau$, we can rewrite (11) as the following approximate inequality.

$$\begin{aligned}
 C_q(t_1) + (t - t_1) - d(\kappa\tau + \xi) &\lesssim C_i(t) \\
 &\lesssim C_q(t_1) + (t - t_1) + d\kappa\tau.
 \end{aligned} \quad (12)$$

Since this holds for all i , we get

$$|C_i(t) - C_j(t)| \lesssim d(2\kappa\tau + \xi),$$

and this holds for all $t \gtrsim t_0 + dt$. ■

Note that relation (11) of the proof yields an exact upper bound for $|C_i(t) - C_j(t)|$ in case the assumption $\mu + \xi \ll \tau$ is invalid. An examination of the proof suggests a simple method for rapidly initializing the clocks, or resynchronizing them if they should go out of synchrony for any reason. Each process sends a message which is relayed to every other process. The procedure can be initiated by any process, and requires less than $2d(\mu + \xi)$ seconds to effect the synchronization, assuming each of the messages has an unpredictable delay less than ξ .

Acknowledgment. The use of timestamps to order operations, and the concept of anomalous behavior are due to Paul Johnson and Robert Thomas.

Received March 1976; revised October 1977

References

- [1] Schwartz, J.T. *Relativity in Illustrations*. New York U. Press, New York, 1962.
- [2] Taylor, E.F., and Wheeler, J.A. *Space-Time Physics*, W.H. Freeman, San Francisco, 1966.
- [3] Lamport, L. The implementation of reliable distributed multiprocess systems. To appear in *Computer Networks*.
- [4] Ellingson, C., and Kulpinski, R.J. Dissemination of system-time. *IEEE Trans. Comm. Com-23*, 5 (May 1973), 605–624.

References

- M. Abadi and L. Lamport. 1991. The existence of refinement mappings. *Theor. Comput. Sci.*, 82(2):253–284. DOI: [10.1016/0304-3975\(91\)90224-P](https://doi.org/10.1016/0304-3975(91)90224-P). 10, 113, 163
- M. Abadi and L. Lamport. 1993. Composing specifications. *ACM Trans. Program. Lang. Syst.*, 15(1): 73–132. DOI: [10.1145/151646.151649](https://doi.org/10.1145/151646.151649). 114
- M. Abadi and L. Lamport. 1994. Open systems in TLA. In J. H. Anderson, D. Peleg, and E. Borowsky (eds.), *Proceedings of the Thirteenth Annual ACM Symposium on Principles of Distributed Computing*, pp. 81–90. ACM. DOI: [10.1145/197917.197960](https://doi.org/10.1145/197917.197960). 114
- M. Abadi and L. Lamport. 1995. Conjoining specifications. *ACM Trans. Program. Lang. Syst.*, 17(3): 507–534. DOI: [10.1145/203095.201069](https://doi.org/10.1145/203095.201069). 114
- J.-R. Abrial. 1996. *The B-Book: Assigning Programs to Meanings*. Cambridge University Press. 115
- J.-R. Abrial. 2010. *Modeling in Event-B*. Cambridge University Press. 115, 127
- S. V. Adve and M. D. Hill. 1990. Weak ordering - A new definition. In *Proceedings of the 17th Annual International Symposium on Computer Architecture*. pp. 2–14, 1990. DOI: [10.1145/325096.325100](https://doi.org/10.1145/325096.325100). 46
- Y. Afek, H. Attiya, D. Dolev, E. Gafni, M. Merritt, and N. Shavit. 1993. Atomic snapshots of shared memory. *J. ACM*, 40(4): 873–890. DOI: [10.1145/153724.153741](https://doi.org/10.1145/153724.153741). 64
- M. Ahamad, G. Neiger, J. E. Burns, P. Kohli, and P. W. Hutto. 1995. Causal memory: Definitions, implementation, and programming. *Distributed Computing*, 9: 37–49. DOI: [10.1007/BF01784241](https://doi.org/10.1007/BF01784241). 53
- B. Alpern and F. B. Schneider. 1985. Defining liveness. *Inf. Process. Lett.*, 21(4): 181–185. DOI: [10.1016/0020-0190\(85\)90056-0](https://doi.org/10.1016/0020-0190(85)90056-0). 104
- P. A. Alsberg and J. D. Day. 1976. A principle for resilient sharing of distributed resources. In *Proceedings of the 2nd International Conference on Software Engineering*, ICSE '76, pp. 562–570. IEEE Computer Society Press. 83, 85
- American Mathematical Society. 2019a. AMS-LaTeX. <http://www.ams.org/publications/authors/tex/amslatex>; last retrieved 10 February 2019. 2, 145
- American Mathematical Society. 2019b. AMS-TeX. <http://www.ams.org/publications/authors/tex/amstex>; last retrieved 10 February 2019. 145

- K. R. Apt. 1981. Ten years of Hoare's logic: A survey—part 1. *ACM Trans. Prog. Lang. Syst.*, 3(4): 431–483, 1981. DOI: [10.1145/357146.357150](https://doi.org/10.1145/357146.357150). 109
- P. C. Attie, N. Francez, and O. Grumberg. 1993. Fairness and hyperfairness in multi-party interactions. *Distributed Computing*, 6(4): 245–254, 1993. DOI: [10.1007/BF02242712](https://doi.org/10.1007/BF02242712). 104
- H. Attiya, F. Ellen, and A. Morrison. 2015. Limitations of highly-available eventually-consistent data stores. In *Proceedings of the 2015 ACM Symposium on Principles of Distributed Computing, PODC '15*, pp. 385–394, New York. ACM. 53
- H. Attiya and J. Welch. 1998. *Distributed Computing: Fundamentals, Simulations, and Advanced Topics*. McGraw-Hill International (UK), London, 1998. 37, 40
- H. Attiya and J. Welch. 2004. *Distributed Computing: Fundamentals, Simulations, and Advanced Topics, Second Edition*. John Wiley & Sons, Hoboken, NJ, 2004. 37
- H. Attiya and J. L. Welch. 1994. Sequential consistency versus linearizability. *ACM Trans. Comput. Syst.*, 12(2): 91–122. DOI: [10.1145/176575.176576](https://doi.org/10.1145/176575.176576). 46
- N. Azmy, S. Merz, and C. Weidenbach. 2018. A machine-checked correctness proof for pastry. *Sci. Comput. Program.*, 158: 64–80. DOI: [10.1016/j.scico.2017.08.003](https://doi.org/10.1016/j.scico.2017.08.003). 127
- R.-J. Back. 1981. On correct refinement of programs. *J. Comput. Syst. Sci.*, 23(1): 49–68, 1981. DOI: [10.1016/0022-0000\(81\)90005-2](https://doi.org/10.1016/0022-0000(81)90005-2). 111
- J. W. Backus, R. J. Beeber, S. Best, R. Goldberg, L. M. Haibt, H. L. Herrick, R. A. Nelson, D. Sayre, P. B. Sheridan, H. Stern. 1957 The Fortran automatic coding system. In *Papers presented at the February 26–28, 1957, Western Joint Computer Conference: Techniques for Reliability*, pp. 188–198. ACM. DOI: [10.1145/1455567.1455599](https://doi.org/10.1145/1455567.1455599). 133
- C. Barrett and C. Tinelli. 2018. Satisfiability modulo theories. In E. M. Clarke, T. A. Henzinger, H. Veith, and R. Bloem (eds.), *Handbook of Model Checking*, pp. 305–343. Springer. 126
- M. Ben-Or. 1983. Another advantage of free choice: Completely asynchronous agreement protocols (extended abstract). In *Proceedings of the Second Annual ACM Symposium on Principles of Distributed Computing, PODC '83*, pp. 27–30. ACM. DOI: [10.1145/800221.806707](https://doi.org/10.1145/800221.806707). 90
- M. Ben-Or and R. El-Yaniv. 2003. Resilient-optimal interactive consistency in constant time. *Distributed Computing*, 16: 249–262. DOI: [10.1007/s00446-002-0083-3](https://doi.org/10.1007/s00446-002-0083-3). 71
- D. J. Bernstein, D. Hopwood, A. Hülsing, T. Lange, R. Niederhagen, L. Papachristodoulou, M. Schneider, P. Schwabe, and Z. Wilcox-O'Hearn. 2015. SPHINCS: Practical stateless hash-based signatures. In *Proceedings EUROCRYPT 2015*, pp. 368–397. DOI: [10.1007/978-3-662-46800-5_15](https://doi.org/10.1007/978-3-662-46800-5_15). 81
- I. Beschastnikh, P. Wang, Y. Brun, and M. D. Ernst. 2016. Debugging distributed systems. *Queue*, 14(2): 50. 53
- K. Birman, A. Schiper, and P. Stephenson. August 1991. Lightweight causal and atomic group multicast. *ACM Trans. Comput. Syst.*, 9(3): 272–314. DOI: [10.1145/128738.128742](https://doi.org/10.1145/128738.128742). 53, 102

- K. Birman. 1986. Isis: A system for fault-tolerant distributed computing. Technical report, Cornell University. 53
- A. D. Birrell, A. Hisgen, C. Jerian, T. Mann, and G. Swart. September 1993. The Echo distributed file system. *Digital Systems Research Center Research Report 111*, 10. DOI: [10.1.1.43.1306](#). 148
- R. L. Bocchino, Jr., V. S. Adve, S. V. Adve, and M. Snir. 2009. Parallel programming must be deterministic by default. In *Proceedings of the First USENIX Conference on Hot Topics in Parallelism*, HotPar'09, pp. 4–10, Berkeley, CA. USENIX Association. 86
- W. J. Bolosky, J. R. Douceur, and J. Howell. 2007. The Farsite project: A retrospective. *Operating Systems Review*, 41(2): 17–26. DOI: [10.1145/1243418.1243422](#). 128
- R. Bonichon, D. Delahaye, and D. Doligez. 2007. Zenon: An extensible automated theorem prover producing checkable proofs. In N. Dershowitz and A. Voronkov (eds.), *14th International Conference Logic for Programming, Artificial Intelligence, and Reasoning (LPAR)*, volume 4790 of *Lecture Notes in Computer Science*, pp. 151–165. Springer. DOI: [10.1007/978-3-540-75560-9_13](#). 126
- F. V. Brasileiro, F. Greve, A. Mostéfaoui, and M. Raynal. 2001. Consensus in one communication step. In *Proceedings of the 6th International Conference on Parallel Computing Technologies*, PaCT '01, pp. 42–50. Springer-Verlag. DOI: [10.1007/3-540-44743-1_4.pdf](#). 93
- E. A. Brewer. July 2000. Towards robust distributed systems (abstract). In *Proceedings of the Nineteenth Annual ACM Symposium on Principles of Distributed Computing*, p. 7. DOI: [10.1145/343477.343502](#). 50
- J. A. Buchmann, E. Dahmen, and A. Hülsing. 2011. XMSS - A practical forward secure signature scheme based on minimal security assumptions. In *Proceedings Workshop on Post-Quantum Cryptography (PQC)*, volume 7071 of *Lecture Notes in Computer Science*, pp. 117–129. Springer. DOI: [10.1007/978-3-642-25405-5_8](#). 81
- M. Burrows. January 2019. Personal communication. 150
- M. Burrows. 2006. The Chubby lock service for loosely-coupled distributed systems. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation*, OSDI '06, pp. 335–350. USENIX Association. 9, 100
- R. M. Burstall. 1974. Program proving as hand simulation with a little induction. In *Information Processing*, pp. 308–312. North-Holland Pub. Co. 105
- V. Bush. 1945. As we may think. *The Atlantic Monthly*, 176(1): pp. 101–108. 131
- C. Cachin, R. Guerraoui, and L. Rodrigues. 2011. *Introduction to Reliable and Secure Distributed Programming (Second Edition)*. Springer, 2011. 54, 69
- C. Cachin, K. Kursawe, F. Petzold, and V. Shoup. August 2001. Secure and efficient asynchronous broadcast protocols. In *Advances in Cryptology - CRYPTO 2001, 21st Annual International Cryptology Conference*, pp. 524–541. 71
- S. Chand, Y. A. Liu, and S. D. Stoller. 2016. Formal verification of multi-Paxos for distributed consensus. In J. S. Fitzgerald, C. L. Heitmeyer, S. Gnesi, and A. Philippou (eds.), *21st*

- International Symposium Formal Methods (FM 2016)*, volume 9995 of *Lecture Notes in Computer Science*, pp. 119–136. DOI: [10.1007/978-3-319-48989-6_8](https://doi.org/10.1007/978-3-319-48989-6_8). 127
- T. D. Chandra, R. Griesemer, and J. Redstone. 2007. Paxos made live: An engineering perspective. In *Proceedings of the 26th ACM Symposium on Principles of Distributed Computing (PODC '07)*, pp. 398–407. ACM. DOI: [10.1145/1281100.1281103](https://doi.org/10.1145/1281100.1281103). 9, 99
- T. D. Chandra and S. Toueg. 1996. Unreliable failure detectors for reliable distributed systems. *J. ACM*, 43(2): 225–267. DOI: [10.1145/226643.226647](https://doi.org/10.1145/226643.226647). 71
- K. M. Chandy and L. Lamport. 1985. Distributed snapshots: Determining global states of distributed systems. *ACM Trans. Comput. Syst.*, 3(1): 63–75. DOI: [10.1145/214451.214456](https://doi.org/10.1145/214451.214456). 7, 48, 59, 64, 143, 164
- M. Chandy. July 2018. Personal communications (email), 16 and 30. 165
- A. Condon and A. J. Hu. 2003. Automatable verification of sequential consistency. *Theory Comput. Syst.*, 36(5): 431–460. DOI: [10.1007/s00224-003-1082-x](https://doi.org/10.1007/s00224-003-1082-x). 46
- D. Cousineau, D. Doligez, L. Lamport, S. Merz, D. Ricketts, and H. Vanzetto. August 2012. TLA+ proofs. In D. Giannakopoulou and D. Méry (eds.), *FM 2012: Formal Methods - 18th International Symposium*, volume 7436 of *Lecture Notes in Computer Science*, pp. 147–154. Springer. 125
- G. DeCandia, D. Hastorun, M. Jampani, G. Kakulapati, A. Lakshman, A. Pilchin, S. Sivasubramanian, P. Vosshall, and W. Vogels. 2007. Dynamo: Amazon’s highly available key-value store. In *SOSP*. DOI: [10.1145/1323293.1294281](https://doi.org/10.1145/1323293.1294281). 53
- D. Didona, R. Guerraoui, J. Wang, and W. Zwaenepoel. 2018. Causal consistency and latency optimality: Friend or foe? *PVLDB*, 11(11): 1618–1632. DOI: [10.14778/3236187.3236210](https://doi.org/10.14778/3236187.3236210). 53
- W. Diffie and M. E. Hellman. 1976. New directions in cryptography. *IEEE Transactions on Information Theory*, 22(6): 644–654. DOI: [10.1109/TIT.1976.1055638](https://doi.org/10.1109/TIT.1976.1055638). 78, 140
- E. W. Dijkstra. 1971. On the reliability of programs. <https://www.cs.utexas.edu/users/EWD/transcriptions/EWD03xx/EWD303.html>. 159
- E. W. Dijkstra. 1965. Solution of a problem in concurrent programming control. *CACM*, 8(9): 569. 4
- E. W. Dijkstra. 1971. Hierarchical ordering of sequential processes. *Acta Inf.*, 1: 115–138. DOI: [10.1007/BF00289519](https://doi.org/10.1007/BF00289519). 30, 44
- E. W. Dijkstra. 1976. *A Discipline of Programming*. Prentice Hall. 111
- D. L. Dill, A. J. Drexler, A. J. Hu, and C. H. Yang. 1992. Protocol verification as a hardware design aid. In *IEEE International Conference Computer Design: VLSI in Computers and Processors*, pp. 522–525. IEEE Computer Society. 123
- D. Dolev and N. Shavit. 1997. Bounded concurrent time-stamping. *SIAM J. Comput.*, 26(2): 418–455. DOI: [10.1137/S0097539790192647](https://doi.org/10.1137/S0097539790192647). 37, 40
- D. Dolev and H. R. Strong. 1982. Polynomial algorithms for multiple processor agreement. In *Proceedings Symposium on Theory of Computing (STOC)*, pp. 401–407. DOI: [10.1145/800070.802215](https://doi.org/10.1145/800070.802215). 76

- D. Dolev and H. R. Strong. 1983. Authenticated algorithms for Byzantine agreement. *SIAM J. Comput.* 12(4): 656–666. DOI: [10.1137/0212045](#). 76
- D. Dolev and A. C. Yao. 1983. On the security of public key protocols. *IEEE Transactions on Information Theory*, 29(2): 198–207. DOI: [10.1109/TIT.1983.1056650](#). 78
- D. Doligez, J. Kriener, L. Lamport, T. Libal, and S. Merz. 2014. Coalescing: Syntactic abstraction for reasoning in first-order modal logics. In C. Benz Müller and J. Otten (eds.), *Automated Reasoning in Quantified Non-Classical Logics*, volume 33 of *EPiC Series in Computing*, pp. 1–16. 126
- J. Du, C. Iorgulescu, A. Roy, and W. Zwaenepoel. 2014. Gentlerain: Cheap and scalable causal consistency with physical clocks. In *Proceedings of the ACM Symposium on Cloud Computing*, SOCC '14, pp. 4: 1–4: 13. ACM. DOI: [10.1145/2670979.2670983](#). 53
- C. Dwork, N. Lynch, and L. Stockmeyer. April 1988. Consensus in the presence of partial synchrony. *J. ACM*, 35(2): 288–323. DOI: [10.1145/42282.42283](#). 94
- C. Dwork and O. Waarts. 1992. Simple and efficient bounded concurrent timestamping or bounded concurrent timestamp systems are comprehensible! In *Proceedings of the 24th Annual ACM Symposium on Theory of Computing*, pp. 655–666. DOI: [10.1145/129712.129776](#). 37, 40
- C. A. Ellis. 1977. Consistency and correctness of duplicate database systems. In *Proceedings of the Sixth ACM Symposium on Operating Systems Principles*, SOSP '77, pp. 67–84. ACM. DOI: [10.1145/1067625.806548](#). 84
- U. Engberg. 1996. *Reasoning in the Temporal Logic of Actions*. BRICS Dissertation Series. 151
- U. Engberg, P. Grønning, and L. Lamport. 1992. Mechanical verification of concurrent systems with TLA. In G. von Bochmann and D. K. Probst (eds.), *Computer Aided Verification, Fourth International Workshop, CAV '92*, volume 663 of *Lecture Notes in Computer Science*, pp. 44–55. Springer. DOI: [10.1007/3-540-56496-9_5](#). 115
- C. J. Fidge. 1988. Timestamps in message-passing systems that preserve the partial ordering. Dissertation, Australian National University. 53
- M. J. Fischer, N. A. Lynch, and M. S. Paterson. April 1985. Impossibility of distributed consensus with one faulty process. *J. ACM*, 32(2): 374–382. DOI: [10.1145/3149.214121](#). 89, 148
- M. Fitzi and J. A. Garay. 2003. Efficient player-optimal protocols for strong and differential consensus. In *Proceedings 22nd ACM Symposium on Principles of Distributed Computing (PODC)*, pp. 211–220. DOI: [10.1145/872035.872066](#). 70
- R. W. Floyd. 1967. Assigning meanings to programs. In *Proceedings Symposium on Applied Mathematics*, volume 19, pp. 19–32. American Mathematical Society. 105
- R. Friedman, A. Mostéfaoui, S. Rajsbaum, and M. Raynal. 2007. Asynchronous agreement and its relation with error-correcting codes. *IEEE Trans. Computers*, 56(7): 865–875. DOI: [10.1109/TC.2007.1043](#). 72
- E. Gafni and L. Lamport. 2003. Disk Paxos. *Distributed Computing*, 16(1): 1–20. DOI: [10.1007/s00446-002-0070-8](#). 9, 100, 161

- J. A. Garay and Y. Moses. 1998. Fully polynomial Byzantine agreement for $n > 3t$ processors in $t + 1$ rounds. 27(1): 247–290. DOI: [10.1137/S0097539794265232](#). 76
- S. J. Garland and J. V. Guttag. 1988. LP: The Larch prover. In *International Conference on Automated Deduction*, pp. 748–749. Springer. DOI: [10.1007/BFb0012879](#). 151
- S. Ghemawat, H. Gobioff, and S.-T. Leung. 2003. The Google file system. In *Proceedings of the Nineteenth ACM Symposium on Operating Systems Principles*, SOSP '03, pp. 29–43. ACM. DOI: [10.1145/1165389.945450](#). 101
- P. B. Gibbons and E. Korach. 1997. Testing shared memories. *SIAM J. Comput.*, 26(4): 1208–1244. DOI: [10.1.1.107.3013](#). 46
- P. B. Gibbons, M. Merritt, and K. Gharachorloo. 1991. Proving sequential consistency of high-performance shared memories (extended abstract). In *SPAA*, pp. 292–303. DOI: [10.1145/113379.113406](#). 46
- S. Gilbert and N. Lynch. 2002. Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services. *SIGACT News*, 33(2): 51–59. DOI: [10.1145/564585.564601](#). 50
- E. Gilkerson and L. Lamport. 2004. On hair color in France. *Annals of Improbable Research*, pp. 18–19. 165
- P. Godefroid and P. Wolper. 1994. A partial approach to model checking. *Inf. Comput.*, 110(2): 305–326. DOI: [10.1006/inco.1994.1035](#). 127
- C. Gray and D. Cheriton. 1989. Leases: An efficient fault-tolerant mechanism for distributed file cache consistency. In *Proceedings of the Twelfth ACM Symposium on Operating Systems Principles*, SOSP '89, pp. 202–210. ACM. DOI: [10.1145/74851.74870](#). 99
- J. Gray and L. Lamport. 2006. Consensus on transaction commit. *ACM Trans. Database Syst.*, 31(1): 133–160, 2006. DOI: [10.1145/1132863.1132867](#). 162
- D. Gries and F. B. Schneider. 1995. Avoiding the undefined by underspecification. In J. van Leeuwen, editor, *Computer Science Today*, volume 1000 of *Lecture Notes in Computer Science*, pp. 366–373. Springer. 116
- Y. Gurevich. 1995. Evolving algebras 1993: Lipari guide. In E. Börger, editor, *Specification and validation methods*, pp. 9–36. Oxford University Press. 127
- C. Hawblitzel, J. Howell, M. Kapritsos, J. R. Lorch, B. Parno, M. L. Roberts, S. T. V. Setty, and B. Zill. 2015. Ironfleet: Proving practical distributed systems correct. In E. L. Miller and S. Hand (eds.), *Proceedings 25th Symposium Operating Systems Principles*, SOSP 2015, pp. 1–17. ACM. DOI: [10.1145/2815400.2815428](#). 128
- C. Hawblitzel, J. Howell, M. Kapritsos, J. R. Lorch, B. Parno, M. L. Roberts, S. T. V. Setty, and B. Zill. 2017. Ironfleet: Proving safety and liveness of practical distributed systems. *Commun. ACM*, 60(7): 83–92. DOI: [10.1145/3068608](#). 128
- M. Herlihy. January 1991. Wait-free synchronization. *ACM Trans. Program. Lang. Syst.*, 13(1): 124–149. DOI: [10.1145/114005.102808](#). 4
- M. Herlihy and J. M. Wing. July 1990. Linearizability: A correctness condition for concurrent objects. *ACM Trans. Prog. Lang. Syst.*, 12(3): 463–492. DOI: [10.1145/78969.78972](#). 5, 71

- A. Hisgen, A. Birrell, T. Mann, M. Schroeder, and G. Swart. September 1989. Availability and consistency tradeoffs in the Echo distributed file system. In *Proceedings of the Second Workshop on Workstation Operating Systems*, pp. 49–54. IEEE Computer Society. DOI: [10.1109/WWOS.1989.109267](https://doi.org/10.1109/WWOS.1989.109267). 96
- G. J. Holzmann. 2003. *The SPIN Model Checker*. Addison-Wesley. 123
- H. Howard, D. Malkhi, and A. Spiegelman. 2016. Flexible Paxos: Quorum intersection revisited. *arXiv preprint arXiv:1608.06696*. 162
- P. Hunt, M. Konar, F. P. Junqueira, and B. Reed. 2010. Zookeeper: Wait-free coordination for internet-scale systems. In *USENIX ATC*. 9, 100
- IBM. 1955. *IBM 705 EDPM Electronic Data Processing Machine*. International Business Machines Corp. 133
- A. Israeli and M. Li. 1993. Bounded time-stamps. *Distributed Computing*, 6(4): 205–209. DOI: [10.1007/BF02242708](https://doi.org/10.1007/BF02242708). 37
- A. Israeli and A. Shaham. 2005. Time and space optimal implementations of atomic multi-writer register. *Inf. Comput.*, 200(1): 62–106. 40
- K. E. Iverson. 1962. A programming language. In *Proceedings of the Spring Joint Computer Conference*, pp. 345–351. ACM. 135
- P. Jayanti. 2005. An optimal multi-writer snapshot algorithm. In *Proceedings of the Thirty-seventh Annual ACM Symposium on Theory of Computing*, pp. 723–732. ACM. DOI: [10.1145/1060590.1060697](https://doi.org/10.1145/1060590.1060697). 64
- J. E. Johnson, D. E. Langworthy, L. Lamport, and F. H. Vogt. 2007. Formal specification of a web services protocol. *J. Log. Algebr. Program.*, 70(1): 34–52. DOI: [10.1016/j.jlap.2006.05.004](https://doi.org/10.1016/j.jlap.2006.05.004). 128
- P. R. Johnson and R. H. Thomas. 1975. The maintenance of duplicate databases. Network Working Group RFC 677. 6, 50, 53, 83, 137
- R. Joshi, L. Lamport, J. Matthews, S. Tasiran, M. R. Tuttle, and Y. Yu. 2003. Checking cache-coherence protocols with TLA⁺. *Formal Methods in System Design*, 22(2): 125–131. DOI: [10.1023/A:1022969405325](https://doi.org/10.1023/A:1022969405325). 123, 127
- R. Koo and S. Toueg. 1987. Checkpointing and rollback-recovery for distributed systems. *IEEE Transactions on software Engineering*, (1): 23–31. DOI: [10.1109/TSE.1987.232562](https://doi.org/10.1109/TSE.1987.232562). 64
- F. Kröger. 1977. LAR: A logic of algorithmic reasoning. *Acta Informatica*, 8: 243–266. DOI: [10.1007/BF00264469](https://doi.org/10.1007/BF00264469). 105
- DEC WRL. October 2018. MultiTitan: Four architecture papers. <http://www.hpl.hp.com/techreports/Compaq-DEC/WRL-87-8.pdf>. 147
- P. B. Ladkin, L. Lamport, B. Olivier, and D. Roegel. 1999. Lazy caching in TLA. *Distributed Computing*, 12(2–3): 151–174. DOI: [10.1007/s004460050063](https://doi.org/10.1007/s004460050063). 155
- A. Lakshman and P. Malik. 2009. Cassandra: A decentralized structured storage system. In *SOSP Workshop on Large Scale Distributed Systems and Middleware (LADIS 2009)*. ACM. 100

- L. Lamport. 1957 Braid theory. *Mathematics Bulletin of the Bronx High School of Science (1957)*, pp. 6,7,9. [132](#)
- L. Lamport. 1987. Distribution. Email message sent to a DEC SRC bulletin board at 12:23:29 PDT on 28 May 87. [163](#)
- L. Lamport. 2019. My writings. [3](#), [29](#), [67](#), [103](#), [132](#), [134](#), [135](#), [141](#), [142](#), [148](#), [150](#), [151](#), [158](#), [165](#)
- L. Lamport. 2018. Personal communication (recorded conversation). [144](#), [145](#)
- L. Lamport. 1970. Comment on Bell's quadratic quotient method for hash coded searching. *Commun. ACM*, 13(9): 573–574. DOI: [10.1145/362736.362765](#). [135](#)
- L. Lamport. 1973. The coordinate method for the parallel execution of DO loops. In *Proceedings 1973 Sagamore Comput. Conf.*
- L. Lamport. 1974. A new solution of Dijkstra's concurrent programming problem. *Commun. ACM*, 17(8): 453–455. DOI: [10.1145/361082.361093](#). [3](#), [29](#), [30](#), [35](#), [136](#)
- L. Lamport. 1974. The parallel execution of DO loops. *Commun. ACM*, 17(2): 83–93. DOI: [10.1145/360827.360844](#). [135](#)
- L. Lamport. 1977. Concurrent reading and writing. *Commun. ACM*, 20(11): 806–811. DOI: [10.1145/359863.359878](#). [30](#), [36](#), [44](#)
- L. Lamport. 1977. Proving the correctness of multiprocess programs. *IEEE Trans. Software Eng.*, 3(2): 125–143. DOI: [10.1109/TSE.1977.229904](#). [30](#), [103](#), [104](#), [105](#), [106](#), [150](#)
- L. Lamport. 1978. The implementation of reliable distributed multiprocess systems. *Computer Networks*, 2: 95–114. DOI: [10.1016/0376-5075\(78\)90045-4](#). [84](#)
- L. Lamport. 1978. Time, clocks, and the ordering of events in a distributed system. *Commun. ACM*, 21(7): 558–565. DOI: [10.1145/359545.359563](#). [6](#), [48](#), [58](#), [103](#), [137](#), [163](#), [164](#)
- L. Lamport. 1979a. Constructing digital signatures from a one way function. *Technical Report CSL—98, Computer Science Laboratory, SRI International*. [68](#), [78](#)
- L. Lamport. 1979b. How to make a multiprocessor computer that correctly executes multiprocess programs. *IEEE Trans. Computers*, 28(9): 690–691. DOI: [10.1109/TC.1979.1675439](#). [5](#), [30](#), [44](#), [143](#), [150](#)
- L. Lamport. 1979c. A new approach to proving the correctness of multiprocess programs. *ACM Trans. Program. Lang. Syst.*, 1(1): 84–97. DOI: [10.1145/357062.357068](#). [103](#)
- L. Lamport. 1980. “Sometime” is sometimes “not never”—on the temporal logic of programs. In P. W. Abrahams, R. J. Lipton, and S. R. Bourne (eds.), *Conference Record of the Seventh Annual ACM Symposium on Principles of Programming Languages*, pp. 174–185. ACM Press. DOI: [10.1145/567446.567463](#). [106](#)
- L. Lamport. 1981. Password authentication with insecure communication. *Commun. ACM*, 24(11): 770–772. DOI: [10.1145/358790.358797](#). [68](#)
- L. Lamport. 1983. What good is temporal logic? In *IFIP Congress*, pp. 657–668. [111](#)
- L. Lamport. 1986a. The mutual exclusion problem: Part I: a theory of interprocess communication. *J. ACM*, 33(2): 313–326. DOI: [10.1145/5383.5384](#). [150](#)
- L. Lamport. 1986b. The mutual exclusion problem: Part II: statement and solutions. *J. ACM*, 33(2): 327–348. DOI: [10.1145/5383.5385](#). [150](#)

- L. Lamport. 1986c. On interprocess communication. Part I: basic formalism. *Distributed Computing*, 1(2): 77–85. DOI: [10.1007/BF01786227](https://doi.org/10.1007/BF01786227). 4, 29, 41, 42, 150
- L. Lamport. 1986d. On interprocess communication. Part II: algorithms. *Distributed Computing*, 1(2): 86–101. DOI: [10.1007/BF01786228](https://doi.org/10.1007/BF01786228). 4, 29, 36, 37, 38, 39, 40, 42, 150
- L. Lamport. 1987. A fast mutual exclusion algorithm. *ACM Trans. Comput. Syst.*, 5(1): 1–11. DOI: [10.1145/7351.7352](https://doi.org/10.1145/7351.7352). 147
- L. Lamport. 1990. *win* and *sin*: Predicate transformers for concurrency. *ACM Trans. Program. Lang. Syst.*, 12(3): 396–428. DOI: [10.1145/78969.78970](https://doi.org/10.1145/78969.78970). 104
- L. Lamport. 1991. The temporal logic of actions. *Research Report 79*, DEC Systems Research Center. 107
- L. Lamport. 1992. Hybrid systems in TLA^+ . In R. L. Grossman, A. Nerode, A. P. Ravn, and H. Rischel (eds.), *Hybrid Systems*, volume 736 of *Lecture Notes in Computer Science*, pp. 77–102. Springer. DOI: [10.1007/3-540-57318-6_25](https://doi.org/10.1007/3-540-57318-6_25). 115
- L. Lamport. June 1993. Verification and specification of concurrent programs. In J. W. de Bakker, W. P. de Roever, and G. Rozenberg (eds.), *A Decade of Concurrency, Reflections and Perspectives*, volume 803 of *Lecture Notes in Computer Science*, pp. 347–374. Springer. 151
- L. Lamport. 1994. $\mathcal{L}^{\mathcal{A}}\text{T}_{\mathcal{E}}\text{X}$ —A Document Preparation System: User’s Guide and Reference Manual, Second Edition. Pearson/Prentice Hall. 11, 145
- L. Lamport. 1994. The temporal logic of actions. *ACM Trans. Program. Lang. Syst.*, 16(3): 872–923. DOI: [10.1145/177492.177726](https://doi.org/10.1145/177492.177726). 107, 109, 151
- L. Lamport. 1995. How to write a proof. *American Mathematical Monthly*, 102(7): 600–608. 125, 152, 154
- L. Lamport. 1998. The part-time parliament. *ACM Trans. Comput. Syst.*, 16(2): 133–169. DOI: [10.1145/279227.279229](https://doi.org/10.1145/279227.279229). 9, 54, 96, 163
- L. Lamport. 1999. Specifying concurrent systems with TLA^+ . *Calculational System Design*. DOI: [10.1007/978-3-540-48153-2_6.pdf](https://doi.org/10.1007/978-3-540-48153-2_6.pdf). 152, 156
- L. Lamport. 2000. Fairness and hyperfairness. *Distributed Computing*, 13(4): 239–245. DOI: [10.1007/PL00008921](https://doi.org/10.1007/PL00008921). 104
- L. Lamport. 2000. How $\mathcal{L}^{\mathcal{A}}\text{T}_{\mathcal{E}}\text{X}$ changed the face of mathematics. E-interview in DMV-Mitteilungen.
- L. Lamport. 2001. Paxos made simple. *SIGACT News*, 32(4): 51–58. 98, 150
- L. Lamport. 2002. *Specifying Systems: The TLA^+ Language and Tools for Hardware and Software Engineers*. Addison-Wesley. 10, 115, 152, 156, 163
- L. Lamport. 2002. A discussion with Leslie Lamport. *IEEE Distributed Systems Online* 3, 8.
- L. Lamport. 2006c. Checking a multithreaded algorithm with ^+CAL . In S. Dolev, editor, *Distributed Computing, 20th International Symposium*, volume 4167 of *Lecture Notes in Computer Science*, pp. 151–163. Springer. DOI: [10.1007/11864219_11](https://doi.org/10.1007/11864219_11). 159

- L. Lamport. 2006b. Fast Paxos. *Distributed Computing*, 19(2): 79–103. DOI: [10.1007/s00446-006-0005-x](https://doi.org/10.1007/s00446-006-0005-x). 161
- L. Lamport. 2006. Measuring celebrity. *Annals of Improbable Research*, pp. 14–15. 165
- L. Lamport. 2009. The PlusCal algorithm language. In M. Leucker and C. Morgan (eds.), *Theoretical Aspects of Computing*, volume 5684 of *Lecture Notes in Computer Science*, pp. 36–60. Springer. DOI: [10.1007/978-3-642-03466-4_2](https://doi.org/10.1007/978-3-642-03466-4_2). 10, 120, 159
- L. Lamport. 2011. Byzantizing Paxos by refinement. In D. Peleg, editor, *Distributed Computing - 25th International Symposium*, volume 6950 of *Lecture Notes in Computer Science*, pp. 211–224. Springer. 127, 162
- L. Lamport. 2012. Buridan’s principle. *Foundations of Physics*, 42(8): 1056–1066. DOI: [10.1007/s10701-012-9647-7](https://doi.org/10.1007/s10701-012-9647-7). 142
- L. Lamport. 2012. How to write a 21st century proof. *Journal of Fixed Point Theory and Applications*, March 6, 2012. DOI: [10.1007/s11784-012-0071-6](https://doi.org/10.1007/s11784-012-0071-6). 163
- L. Lamport. 2014. TLA⁺²: A preliminary guide.
- L. Lamport. 2015. The TLA⁺ hyperbook. 115
- L. Lamport. 2015a. Turing lecture: The computer science of concurrency: the early years. *Commun. ACM*, 58(6): 71–76. DOI: [10.1145/2771951](https://doi.org/10.1145/2771951).
- L. Lamport. 2018. The TLA⁺ video course. 115
- L. Lamport. August 2018b. Personal communication (email). 131, 132
- L. Lamport. 2018c. Personal communication (email), 9 October 2018. 152
- L. Lamport. 2019. TLA+ tools. <http://lamport.azurewebsites.net/tla/tools.html>; last retrieved 11 February 2019. 160
- L. Lamport and R. Levin. 2016. Lamport, Leslie. Oral history, part 1. <https://www.computerhistory.org/collections/catalog/102717182>. 2, 132, 133, 134, 135, 136, 137, 138, 139, 141, 142, 146, 148, 149, 159, 163
- L. Lamport and R. Levin. 2016. Lamport, Leslie. Oral history, part 2. <https://www.computerhistory.org/collections/catalog/102717246>. 169
- L. Lamport, D. Malkhi, and L. Zhou. 2009a. Vertical Paxos and primary-backup replication. In S. Tirthapura and L. Alvisi (eds.), *Proceedings of the 28th Annual ACM Symposium on Principles of Distributed Computing*, pp. 312–313. ACM. DOI: [10.1145/1582716.1582783](https://doi.org/10.1145/1582716.1582783). 9, 101, 162
- L. Lamport, D. Malkhi, and L. Zhou. April 2009b. Stoppable Paxos. 162
- L. Lamport, D. Malkhi, and L. Zhou. 2010. Reconfiguring a state machine. *SIGACT News*, 41(1): 63–73. DOI: [10.1145/1753171.1753191](https://doi.org/10.1145/1753171.1753191). 9, 102, 162
- L. Lamport and M. Massa. 2004. Cheap Paxos. In *2004 International Conference on Dependable Systems and Networks*, pp. 307–314. IEEE Computer Society. 9, 101, 161
- L. Lamport, J. Matthews, M. R. Tuttle, and Y. Yu. 2002. Specifying and verifying systems with TLA⁺. In G. Muller and E. Jul (eds.), *Proceedings of the 10th ACM SIGOPS European Workshop*, pp. 45–48. ACM. 153, 154, 157

- L. Lamport and P. M. Melliar-Smith. 1985. Synchronizing clocks in the presence of faults. *J. ACM*, 32(1): 52–78. DOI: [10.1145/2455.2457](https://doi.org/10.1145/2455.2457). 8, 76
- L. Lamport and S. Merz. 2017. Auxiliary variables in TLA+. *CoRR*, abs/1703.05121. 113
- L. Lamport and L. C. Paulson. 1999. Should your specification language be typed? *ACM Trans. Program. Lang. Syst.*, 21(3): 502–526. DOI: [10.1145/319301.319317](https://doi.org/10.1145/319301.319317). 116
- L. Lamport, R. E. Shostak, and M. C. Pease. 1982. The Byzantine generals problem. *ACM Trans. Program. Lang. Syst.*, 4(3): 382–401. DOI: [10.1145/357172.357176](https://doi.org/10.1145/357172.357176). 7, 8, 67, 68, 69, 70, 72, 73, 75, 76, 105, 141, 163
- L. Lamport. 1972. *The Analytic Cauchy Problem with Singular Data*. Ph.D. thesis, Brandeis University.
- B. W. Lampson. 1996. How to build a highly available system using consensus. In *Workshop on Distributed Algorithms (WDAG)*, pp. 1–17. Springer. DOI: [10.1007/3-540-61769-8_1](https://doi.org/10.1007/3-540-61769-8_1). 9
- B. W. Lampson. 2018. Personal communication (email). 139, 140, 149, 150, 167
- B. W. Lampson and H. E. Sturgis. 1979. Crash recovery in a distributed storage system. Unpublished manuscript. 162, 167
- D. Langworthy. 2005. Personal communication (email). 158
- The L^AT_EX Project. 2019. <https://www.latex-project.org/>; last retrieved 6 May 2019. 145
- E. K. Lee and C. A. Thekkath. 1996. Petal: Distributed virtual disks. In *ACM SIGPLAN Notices*, volume 31, pp. 84–92. ACM. DOI: [10.1145/248209.237157](https://doi.org/10.1145/248209.237157). 100, 149
- K. R. M. Leino. 2010. Dafny: An automatic program verifier for functional correctness. In E. M. Clarke and A. Voronkov (eds.), *16th International Conference Logic for Programming, Artificial Intelligence, and Reasoning*, volume 6355 of *Lecture Notes in Computer Science*, pp. 348–370. Springer. DOI: [10.1007/978-3-642-17511-4_20](https://doi.org/10.1007/978-3-642-17511-4_20). 128
- O. Lichtenstein and A. Pnueli. 1985. Checking that finite state concurrent programs satisfy their linear specification. In M. S. V. Deusen, Z. Galil, and B. K. Reid (eds.), *Proceedings Twelfth Ann. ACM Symposium Princ. Prog. Lang. (POPL 1985)*, pp. 97–107. ACM. DOI: [10.1145/318593.318622](https://doi.org/10.1145/318593.318622). 124
- R. Lipton and J. Sandberg. 1988. PRAM: A scalable shared memory. Technical Report CS-TR-180-88, Computer Science Department, Princeton University. 46
- B. Liskov and J. Cowling. 2012. Viewstamped replication revisited. 2012. 54
- B. Liskov and R. Ladin. 1986. Highly available distributed services and fault-tolerant distributed garbage collection. In *Proceedings of the Fifth Annual ACM Symposium on Principles of Distributed Computing, PODC '86*, pp. 29–39. ACM. DOI: [10.1145/10590.10593](https://doi.org/10.1145/10590.10593). 53
- W. Lloyd, M. J. Freedman, M. Kaminsky, and D. G. Andersen. 2011. Don't settle for eventual: Scalable causal consistency for wide-area storage with COPS. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles, SOSP '11*, pp. 401–416, New York, NY. ACM. DOI: [10.1145/2043556.2043593](https://doi.org/10.1145/2043556.2043593). 53

- W. Lloyd, M. J. Freedman, M. Kaminsky, and D. G. Andersen. 2013. Stronger semantics for low-latency geo-replicated storage. In *10th USENIX Symposium on Networked Systems Design and Implementation (NSDI 13)*, pp. 313–328. USENIX. 53
- N. A. Lynch. 1996. *Distributed Algorithms*. Morgan Kaufmann, San Francisco. 76
- D. Malkhi. 2018. Personal communication (recorded conversatino). 168
- F. Mattern. 1989. Virtual time and global states of distributed systems. In M. Cosnard et al. (eds.), *Proceedings of the International Workshop on Parallel Algorithms*. Elsevier. 53, 64
- K. McMillan. 1993. *Symbolic Model Checking*. Kluwer Academic Publishers. 123
- R. C. Merkle. 1980. Protocols for public key cryptosystems. In *Proceedings IEEE Symposium on Security and Privacy*, pp. 122–134. DOI: 10.1109/SP.1980.10006. 80
- S. Merz. 1999. A more complete TLA. In J. Wing, J. Woodcock, and J. Davies (eds.), *FM'99: World Congress on Formal Methods*, volume 1709 of *Lecture Notes in Computer Science*, pp. 1226–1244. Springer. DOI: 10.1007/3-540-48118-4_15. 110
- S. Merz. 2008. The specification language TLA⁺. In D. Bjørner and M. C. Henson (eds.), *Logics of Specification Languages*, Monographs in Theoretical Computer Science, pp. 401–451. Springer, Berlin-Heidelberg. DOI: 10.1007/978-3-540-74107-7_8. 115, 160
- S. Merz. July 2018. Personal communication (email). 167
- Microsoft Research. 2019. *TLA⁺ Proof System*. <https://tla.msr-inria.inria.fr/tlaps/content/Home.html>; last retrieved 11 February 2019. 159
- D. Milojeic. 2002. A discussion with Leslie Lamport. IEEE Distributed Systems Online 3, 8. <https://www.microsoft.com/en-us/research/publication/discussion-leslie-lamport/>; last retrieved 6 May 2019. 161
- H. Minkowski. 2017. Translation:space and time—wikisource. [Online; accessed 15 March 2019]. 138
- C. Morgan. 1990 *Programming from specifications*. Prentice Hall. 111
- A. Mullery. December 1971. The distributed control of multiple copies of data. Technical Report RC 3642, IBM, Yorktown Heights, New York. 83
- M. Naor and M. Yung. 1989. Universal one-way hash functions and their cryptographic applications. In *Proceedings Symposium on Theory of Computing (STOC)*, pp. 33–43. DOI: 10.1145/73007.73011. 80
- National Institute of Standards and Technology. 2018. Post-quantum cryptography. Available at <https://csrc.nist.gov/projects/post-quantum-cryptography/>. 81
- C. Newcombe. 2012. Post on TLA+ discussion group. [https://groups.google.com/forum/#!searchin/tlaplus/professional\\$20career/tlaplus/ZJCi-UF31fc/Mawwwi6U1CYJ](https://groups.google.com/forum/#!searchin/tlaplus/professional$20career/tlaplus/ZJCi-UF31fc/Mawwwi6U1CYJ) 160
- C. Newcombe, T. Rath, F. Zhang, B. Munteanu, M. Brooker, and M. Deardeuff. 2015. How Amazon web services uses formal methods. *CACM*, 58(4): 66–73. DOI: 10.1145/2699417. 128, 161

- M. J. K. Nielsen. Titan system manual. <http://www.hpl.hp.com/techreports/Compaq-DEC/WRL-86-1.pdf>, last retrieved 15 October 2018. 147
- B. M. Oki and B. H. Liskov. 1988. Viewstamped replication: A new primary copy method to support highly-available distributed systems. In *Proceedings of the Seventh Annual ACM Symposium on Principles of Distributed Computing*, pp. 8–17, ACM. DOI: [10.1145/62546.62549](https://doi.org/10.1145/62546.62549). 99, 149
- D. Ongaro and J. Ousterhout. 2014. In search of an understandable consensus algorithm. In *Proceedings USENIX Annual Technical Conference ATC*. 54, 100
- S. S. Owicki and D. Gries. 1976. Verifying properties of parallel programs: An axiomatic approach. *Commun. ACM*, 19(5): 279–285. DOI: [10.1145/360051.360224](https://doi.org/10.1145/360051.360224). 30, 44
- S. S. Owicki and L. Lamport. Proving liveness properties of concurrent programs. *ACM Trans. Program. Lang. Syst.*, 4(3): 455–495, 1982. DOI: [10.1145/357172.357178](https://doi.org/10.1145/357172.357178). 106, 107, 108
- D. S. Parker, G. J. Popek, G. Rudisin, A. Stoughton, B. J. Walker, E. Walton, J. M. Chow, D. Edwards, S. Kiser, and C. Kline. 1983. Detection of mutual inconsistency in distributed systems. *IEEE Transactions on Software Engineering*, (3): 240–247. DOI: [10.1109/TSE.1983.236733](https://doi.org/10.1109/TSE.1983.236733). 53
- L. C. Paulson. 1994. *Isabelle: A Generic Theorem Prover*, volume 828 of Lecture Notes in Computer Science. Springer Verlag, Berlin, Heidelberg. See also the Isabelle home page at <http://isabelle.in.tum.de/>. 126
- M. C. Pease, R. E. Shostak, and L. Lamport. 1980. Reaching agreement in the presence of faults. *J. ACM*, 27(2): 228–234. DOI: [10.1145/322186.322188](https://doi.org/10.1145/322186.322188). 7, 8, 67, 68, 75, 141, 163
- A. Pnueli. 1977. The temporal logic of programs. In *Proceedings 18th Annual Symposium on the Foundations of Computer Science*, pp. 46–57. IEEE. DOI: [10.1109/SFCS.1977.32.105](https://doi.org/10.1109/SFCS.1977.32.105)
- R. Prakash and M. Singhal. 1996. Low-cost checkpointing and failure recovery in mobile computing systems. *IEEE Transactions on Parallel and Distributed Systems*, 7(10):1035–1048. DOI: [10.1109/71.539735](https://doi.org/10.1109/71.539735). 64
- A. N. Prior. 1967. *Past, Present and Future*. Clarendon Press, Oxford, U.K. 105
- Riak. Riak KV. <http://basho.com/products/riak-kv>. 53
- R. L. Rivest, A. Shamir, and L. M. Adleman. 1978. A method for obtaining digital signatures and public-key cryptosystems. *Commun. ACM*, 21(2): 120–126. DOI: [10.1145/359340.359342](https://doi.org/10.1145/359340.359342). 80
- T. R. Rodeheffer. 2018. Personal communication (recorded conversation and follow-up email), 9 October 2018 and 11 February 2019. 167
- J. Rompel. 1990. One-way functions are necessary and sufficient for secure signatures. In *Proceedings Symposium on Theory of Computing (STOC)*, pp. 387–394. DOI: [10.1145/100216.100269](https://doi.org/10.1145/100216.100269). 80
- J. H. Saltzer, D. P. Reed, and D. D. Clark. November 1984. End-to-end arguments in system design. *ACM Trans. Comput. Syst.*, 2(4): 277–288. DOI: [10.1145/357401.357402](https://doi.org/10.1145/357401.357402). 86

- M. Satyanarayanan, J. J. Kistler, P. Kumar, M. E. Okasaki, E. H. Siegel, and D. C. Steere. 1990. Coda: A highly available file system for a distributed workstation environment. *IEEE Transactions on Computers*, 39(4): 447–459. DOI: [10.1109/12.54838](https://doi.org/10.1109/12.54838). 53
- F. B. Schmuck. 1988. The use of efficient broadcast protocols in asynchronous distributed systems. Ph.D. thesis, Cornell University. 53
- F. Schneider. August 2018. Personal communication (email). 166
- F. B. Schneider. 1982. Synchronization in distributed programs. *ACM Trans. Program. Lang. Syst.*, 4(2): 125–148. DOI: [10.1145/357162.357163](https://doi.org/10.1145/357162.357163). 139
- F. B. Schneider. December 1990. Implementing fault-tolerant services using the state machine approach: A tutorial. *ACM Comput. Surv.*, 22(4): 299–319. DOI: [10.1145/98163.98167](https://doi.org/10.1145/98163.98167). 86, 139
- R. L. Schwartz, P. M. Melliar-Smith, and F. H. Vogt. 1984. An interval-based temporal logic. In E. M. Clarke and D. Kozen (eds.), *In Proceedings Logics of Programs*, volume 164 of *Lecture Notes in Computer Science*, pp. 443–457. Springer. 106
- R. Shostak. July 2018. Personal communication (email), 6–8. 140
- A. K. Singh, J. H. Anderson, and M. G. Gouda. 1994. The elusive atomic register. *J. ACM*, 41(2): 311–339. 37
- J. M. Spivey. 1992. *The Z Notation: A Reference Manual*. International Series in Computer Science. Prentice Hall, 2nd edition. 115
- M. Suda and C. Weidenbach. 2012. A PLTL-prover based on labelled superposition with partial model guidance. In B. Gramlich, D. Miller, and U. Sattler (eds.), *6th International Joint Conference Automated Reasoning (IJCAR 2012)*, volume 7364 of *LNCS*, pp. 537–543. Springer. DOI: [10.1007/978-3-642-31365-3_42](https://doi.org/10.1007/978-3-642-31365-3_42). 126
- D. Terry, M. Theimer, K. Petersen, A. Demers, M. Spreitzer, and C. Hauser. 1995. Managing update conflicts in Bayou, a weakly connected replicated storage system. In *In Proceedings of the Fifteenth ACM Symposium on Operating Systems Principles*, pp. 172–183. DOI: [10.1145/224057.224070](https://doi.org/10.1145/224057.224070). 53
- A. Valmari. June 1990. A stubborn attack on state explosion. In *2nd International Wsh. Computer Aided Verification*, volume 531 of *LNCS*, pp. 156–165, Rutgers. Springer. DOI: [10.1007/BF00709154](https://doi.org/10.1007/BF00709154). 127
- R. van Renesse and F. B. Schneider. 2004. Chain replication for supporting high throughput and availability. In *Proceedings of the 6th Conference on Symposium on Operating Systems Design & Implementation*. USENIX Association. 84, 101
- K. Vidyasankar. 1988. Converting Lamport’s regular register to atomic register. *Inf. Process. Lett.*, 28(6): 287–290. DOI: [10.1016/0020-0190\(88\)90175-5](https://doi.org/10.1016/0020-0190(88)90175-5). 39
- P. M. B. Vitányi and B. Awerbuch. 1986. Atomic shared register access by asynchronous hardware (detailed abstract). In *27th Annual Symposium on Foundations of Computer Science*, pp. 233–243. 40
- W. Vogels. 2009. Eventually consistent. *Commun. ACM*, 52(1): 40–44. DOI: [10.1145/1435417.1435432](https://doi.org/10.1145/1435417.1435432). 83

- P. Voldemort. Voldemort. <https://www.project-voldemort.com/voldemort>. 53
- S. A. Weil, S. A. Brandt, E. L. Miller, D. D. E. Long, and C. Maltzahn. 2006. Ceph: A scalable, high-performance distributed file system. pp. 307–320. 100
- J. H. Wensley, L. Lamport, J. Goldberg, M. W. Green, K. N. Levitt, P. M. Melliar-Smith, R. E. Shostak, and C. B. Weinstock. 1978. Sift: Design and analysis of a fault-tolerant computer for aircraft control. *Proceedings of the IEEE*, 66(10): 1240–1255. DOI: [10.1109/PROC.1978.11114](https://doi.org/10.1109/PROC.1978.11114). 140, 164
- J. H. Wensley, L. Lamport, J. Goldberg, M. W. Green, K. N. Levitt, P. M. Melliar-Smith, R. E. Shostak, and C. B. Weinstock. 1978. Synchronizing clocks in the presence of faults. *Proceedings of the IEEE*, 66(10): 1240–1255. 67
- Wikipedia contributors. 2018a. Bravo (software)—Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Bravo_\(software\)](https://en.wikipedia.org/w/index.php?title=Bravo_(software)). [Online; accessed 25 February 2019]. 143
- Wikipedia contributors. 2018b. Eckhard Pfeiffer—Wikipedia, the free encyclopedia. https://en.wikipedia.org/w/index.php?title=Eckhard_Pfeiffer, 2018. [Online; accessed 26 February 2019]. 155
- Wikipedia contributors. 2018c. Scribe (markup language)—Wikipedia, the free encyclopedia. [https://en.wikipedia.org/w/index.php?title=Scribe_\(markup_language\)](https://en.wikipedia.org/w/index.php?title=Scribe_(markup_language)), 2018. [Online; accessed 25 February 2019]. 143
- Wikipedia contributors. 2018d. UNIVAC—Wikipedia, the free encyclopedia. <https://en.wikipedia.org/w/index.php?title=UNIVAC>, 2018. [Online; accessed 24 February 2019]. 132
- Wikipedia contributors. 2019a 1973 oil crisis—Wikipedia, the free encyclopedia. https://en.wikipedia.org/w/index.php?title=1973_oil_crisis. [Online; accessed 25 February 2019]. 140
- Wikipedia contributors. 2019b. DEC Alpha—Wikipedia, the free encyclopedia. https://en.wikipedia.org/w/index.php?title=DEC_Alpha. [Online; accessed 25 February 2019]. 152
- Wikipedia contributors. 2019c. Illiac—Wikipedia, the free encyclopedia. <https://en.wikipedia.org/w/index.php?title=ILLIAC>. [Online; accessed 25 February 2019]. 135
- Wikipedia contributors. 2019d TeX—Wikipedia, the free encyclopedia. <https://en.wikipedia.org/w/index.php?title=TeX>. [Online; accessed 25 February 2019]. 144, 145
- Y. Yu. August 2018. Personal communication (recorded conversation). 153, 166
- M. Zawirski, N. Preguiça, S. Duarte, A. Bieniusa, V. Balesgas, and M. Shapiro. 2015. Write fast, read in the past: Causal consistency for client-side applications. In *Proceedings of the 16th Annual Middleware Conference*, pp. 75–87. ACM. DOI: [10.1145/2814576.2814733](https://doi.org/10.1145/2814576.2814733). 53
- L. Zhou. 2 August 2018. Personal communication (email). 165
- G. Ziegler. January 2000. How $\mathbb{L}\text{T}_{\mathbb{E}}\text{X}$ changed the face of mathematics: An e-interview with Leslie Lamport, the author of $\mathbb{L}\text{T}_{\mathbb{E}}\text{X}$. *Mitteilungen der Deutschen Mathematiker-Vereinigung*, pp. 49–51. Personal communication (email), 2 August 2018. 145

Index

Footnotes are indicated by an ‘n.’

- Abadi, Martin, 10, 114
- Abortion safe behavior, 254
- Acceptors, 96
- Active replication, 85
- Admissible executions, 31–32, 34
- Agreement. *See also* Consensus
 - problem, 7–8, 68–81
 - property, 75, 86
 - protocols, 161–162
- Airplane flight control system, 7, 67, 140
- Almost-Asynchrony benign system model, 87
- Alpha architecture, 151–152
- AmazonWeb Services, 128
- Anomalous behavior, 189–191
- Antoniadis, Karolos, 343
- APL language, 135n
- Arbiter problem, 142, 243
- Arbiters, 20–22
- Armstrong, Susan, 170
- Array processing computer, 135
- Assembly language programming, 133
- Asynchronous consensus protocols, 89–96
- Asynchronous systems, 54, 59, 69
- Asynchrony benign system model, 87
- Atomicity, 4–5, 29–30, 38, 120–121
- Attiya, Hagit, 343–344
- Auxiliary variables, 113
- Averaging time values, 77
- Avionics control system, 7, 67, 140
- Bakery algorithm, 1–5, 15–17, 29–36, 136–137
- Ballots, 96, 282–287
- Basic protocol, 282, 290–291, 307–311
- Batson, Brannon, 160
- Behaviors, 107
- Ben-Or consensus protocol, 90–93
- Boolean function, 265
- Branching-time temporal logics, 106
- Brandeis University, 134
- Bravo system, 143
- Breadth-first searches, 124
- Broadcast primitives, 69
- Broadcasts, 55, 69–74, 88–89, 92, 308
- Bronx High School of Science, 131
- Bureaucrats, 298–300
- Bush, Vannevar, 131
- Bussan, Mimi, 170
- Byzantine agreement, 7–8, 68–81
- Byzantine consensus, 70
- Byzantine generals problem, 67–74, 203
- C-Agreement property, 89
- C-Completion property, 88–89
- C-Validity property, 89

- Cachin, Christian, 40, 344
- CAP theorem, 50
- Castro-Liskov algorithm, 162
- Causality, 49–52
- Chain replication, 84
- Chandy, Mani, 6, 143, 164–165
- Cheap Paxos protocol, 101, 161
- CISC (Complex Instruction Set Computers), 151–152
- Clocks, 5–6, 51, 76–77, 182–185, 189–192
- Closed systems, 114
- Coalescing transformations, 126
- Coherent cache behavior, 5
- Commanders, 69
- Commit protocols, 307
- Communication links and paths, 54–55, 217–221
- Communication variables, 240
- Compaq Computer Company, 155
- COMPASS years (1970–1977), 134–139
- Completion property, 86
- Complex Instruction Set Computers (CISC), 151–152
- Composition, 113–114
- Computer-based document production systems, 143–144
- Computer-building project, 132
- Computer Science Lab (CSL), 140
- Concurrency, 13–25, 41–43, 134, 143, 150, 156, 230, 232
- Consensus, 71. *See also* Agreement protocols, 88–89
- Consistency, 39–40, 149
- Consistent cuts, 24
- Consolidated Edison, 133
- ConstantCheck procedure, 56–57
- Constants, 107
- Control flow, 120
- Correct processes, 68–70, 86
- Counters, 137
- Crash failures, 87–88
- Critical sections, 3, 14, 16, 20–21, 30–36, 174–177, 198–199, 235, 240, 249–259
- Cryptography, 77–79
- CSL (Computer Science Lab), 140
- Cycles, 265
- Deadlock Freedom property, 251, 263, 271
- Deadlocks, 30–32, 64, 250–251, 268
- DEC/Compaq years (1985–2001), 146–155
- Decree-ordering property, 297
- Diffie, W., 77–78, 140–141
- Digital signatures, 71–74, 77–81
- Dijkstra, Edsger, 2–4, 14–15, 18, 24–25, 143, 159
- Dining philosophers problem, 3
- Dishonest legislators, 302–303
- Disk Paxos protocol, 100, 161
- Distributed algorithms, introduction, 23–24
- Distributed global state, 58–64
- Distributed snapshots, 58–63, 143
- Distributed state machine abstraction, 53–58
- Dwork, Cynthia, 94
- Dwork, Lynch, and Stockmeyer consensus protocols, 94–96
- Dynamic reconfiguration, 100–102
- Early employment years, 133–134
- Echo Distributed File System, 148
- Eclipse platform, 160
- Edsger W. Dijkstra Prize in Distributed Computing, 8, 141
- Education, 133–134
- Effectively nonconcurrent operations, 238
- Effectively precedes operations, 237
- Either primitive, 120
- Election prediction by computer, 132
- Elementary operation executions, 234
- Elements of cycles, 265
- Ellis, Clarence, 84
- Engberg, Urban, 151
- EV6 multiprocessor, 152
- EV7 multiprocessor, 153–154
- EV8 multiprocessor, 154
- Events, 19–20, 23–24, 47–52, 229–231
- Eventually consistent databases, 83

- Existential forgery, 78
- Exit conditions, 32, 250–251, 253
- Exponential information gathering, 76
- Fail-safe property, 256
- Fail-stop failures, 88
- Failure detection oracles, 88
- Failure operation, 256
- Fair keywords, 18
- Fairness requirements, 251–253
- Farsite project, 128
- Fast Paxos protocol, 161
- Fast reads, 304–306
- Fault tolerance, 24–25, 58, 67–68, 137, 140
- FCFS (first-come-first-served) condition, 252–253, 258, 269–273
- FIFO communication links, 54
- FIFO queues, 17–22, 44, 117–119, 122, 125
- Firefly project, 147
- First-come-first-served (FCFS) condition, 252–253, 258, 269–273
- Fischer, Michael, 253
- Flexible Paxos protocol, 162n
- Floyd-Hoare style of reasoning, 128
- Floyd's method, 105
- FLP result, 148n
- Forged digital signatures, 79
- Forget actions, 312
- Formal specification and verification, 9–10, 103–105
 - impact, 127–129
 - PlusCal language, 119–122
 - TLA, 105–114
 - TLA+ language, 115–119
 - TLAPS system, 125–127
 - TLC, 123–125
- FORTTRAN compilers, 134–136
- Foxboro Computers, 136n
- Frangipani system, 149
- Franken consensus, 94–95
- Gafni, Eli, 161
- Garbage collection paper, 143
- Gharachorloo, Kourosh, 152
- Gilkerson, Ellen, 170
- Glitch, The, 142
- Global state, 47–49, 58–64
- Goldberg, Jack, 141
- Google File System, 101
- Graduate education, 134
- Grain of atomicity, 120
- Grant service operations, 239
- Gray, Jim, 141, 162
- Guerraoui, Rachid, 344
- Guibas, Leo, 148
- Guindon, Dick, 162
- Happened before relations, 5, 181
- Hellman, M. E., 77–78, 140–141
- Hexagon argument, 8
- Hiding internal state, 111–113
- Hierarchical decomposition, 154
- Higher-level operation, 42
- Higher-level views, 233–234
- History variables, 10, 113
- Honest mistakes, 302–303
- Honors, 163–165
- IBM 705 computer, 133
- ILLIAC IV computer, 135–136
- Impossibility results, 72, 206–209
- Information hiding, 112
- Initial declaration, 250
- Initial states, 86
- Instantiation, 118–119
- Interactive consistency, 69–71, 73
- Interprocess communication, 234–239
- Interrupts, 136n
- Invariant predicates, 19
- IronFleet project, 128
- Jean-Claude Laprie Award in Dependable Computing, 8
- Johnson, James, 158
- Johnson, Paul R., 50, 53, 137–138
- Johnson, Samuel, 13
- Keidar, Idit, 344

- Key generation, 78–80
- Knuth, Donald, 144
- Kuppe, Markus, 160
- Ladkin, Peter, 148
- Lamport, Leslie, biography
 - collegial influences, 165–169
 - COMPASS years (1970–1977), 134–139
 - DEC/Compaq years (1985–2001), 146–155
 - early years, 131–132
 - education and early employment, 133–134
 - honors, 163–165
 - Microsoft years (2001–), 155–163
 - photo, 170–172
 - SRI years (1977–1985), 139–146
- Lamport signatures, 68, 79–81
- Lamport timestamps, 5, 51–52
- Lampson, Butler, 139n, 143, 149n, 162n, 167
- Langworthy, David, 158
- Latest update wins replication technique, 83
- L^AT_EX system, 11, 143–145
- Leaders, 56–57, 96–98
- Learning the law process, 300–302
- Lee, Ed, 149
- Legislator selection, 303–304
- Levin, Roy, 155, 344
- LICS Test-of-Time Award, 10
- Lieutenants, 69
- Linear-time temporal logics, 106
- Linearizability, 5
- Livelock freedom, 14
- Liveness properties, 14, 104–105, 124, 127
- Lockout Freedom Property, 251, 256
- Lockouts, 31–32, 268
- Logic TLA, 107–110
- Logical clocks, 5–6, 51, 182–185
- Logical time, 48–58
- Logical timestamps, 5
- Long ledgers, 298
- Lord, Alison, 132
- Lynch, Nancy, 94
- Majority decisions, 76
- Malfunction execution, 256
- Malfunctioning behavior, 255
- Malicious intelligence, 225
- Malkhi, Dahlia, 161, 168, 343
- Mann, Tim, 149
- Marked graphs, 21–22
- Marker messages, 60–62
- Marlboro College, 134
- Marzullo, Keith, 149, 159, 279
- Massa, Mike, 161
- Massachusetts Computer Associates (COMPASS), 134–139
- Mathematical results, 282–287
- Mathematics, graduate education in, 134
- May affect relations, 103
- Median time values, 77
- Melliard-Smith, Michael, 8
- Memory Barrier (MB) instructions, 16–17
- Memory consistency models, 5. *See also* Shared memory
- Merkle signature scheme, 80–81
- Merz, Stephan, 167, 344–345
- Metaphors, 2
- Microsoft years (2001–), 155–163
- MicroVAX line, 147
- Milner, R., 25
- Minimum delay of messages, 191
- Minkowski, H., 138
- Missing communication paths, 217–221
- MIT, 133
- MITRE Corporation, 134
- Model checkers, 123–127
- Modules, 116–119
- Molnar, Charles, 142
- Monotonicity condition, 300–302
- Moore, J., 152
- Multi-Paxos protocol, 9, 99
- Multidecree parliament, 294–297
- Multiple-reader variables, 242–243
- MultiTitan multiprocessor, 147
- Murphi model checker, 123
- Mutual exclusion
 - Bakery algorithm, 30–32

- MultiTitan multiprocessor, 147
- problem, 14–15
- vs. producer-consumer synchronization, 20–21
- proof, 16–17
- real solution, 15–16
- shared memory, 29–30
- simplification of, 136
- work on, 150
- Mutual Exclusion Problem: Statement and Solutions
 - abstract, 247
 - basic requirements, 249–251
 - conclusion, 273–275
 - failure, 254–258
 - fairness requirements, 251–253, 256
 - FCFS solutions, 269–273
 - introduction, 248
 - One-Bit Algorithm, 260–264
 - premature termination, 253–254
 - problem, 249–258
 - solutions, 258–273
 - Three-Bit Algorithm, 266–269
- Mutual Exclusion Problem: Theory of Interprocess Communication
 - abstract, 227–228
 - assumptions, 243–244
 - conclusion, 244
 - higher-level views, 233–234
 - interprocess communication, 234–239
 - introduction, 228
 - model, 228–234
 - multiple-reader variables, 242–243
 - physical considerations, 229–231
 - processes, 239–241
 - system executions, 231–233
- Mutual exclusion property, 249, 268
- Mutual exclusion protocol, 259–260
- N*-Bit Algorithm, 270–273
- N*-process systems, 21
- National Academy of Engineering honor, 164
- National Institute of Standards and Technology (NIST), 81
- Neighbors, 217
- New-old inversion behavior, 37
- Newcombe, Chris, 128, 160
- Next-time operator, 108
- NIST (National Institute of Standards and Technology), 81
- Nodes, 48
- Nonconcurrent operations, 238
- Noncritical sections, 174, 235, 240, 256–257
- Normal operation execution, 256
- Oki, Brian, 149
- One-Bit Algorithm, 260–264
- One-time signatures, 68, 80–81
- One-way functions, 79
- Open systems, 114
- Operation executions, 16–17, 22–23, 41–42, 231–233
- Operators, 116–118
- Ordered cycles, 265
- Ordering events totally, 185–189
- Ordering of decrees, 296–297
- Owicki, Susan S., 106
- p*-regular graphs, 217
- Palais, Richard, 142, 171
- Parables, 2
- Parallel composition, 113
- Parallel computing. *See* Concurrency
- Parallelism, 135–136
- Partial correctness property, 10
- Partial event ordering, 180–182
- Passed decrees, 296
- Passive replication, 85
- Password authentication, 68
- Paxos algorithm, 9, 148–150, 161–162, 164
- Paxos protocol, 96–100
- PC algorithm, 18–19
- Pease, Marshall, 7–8, 140–141
- Perfect communication links, 54
- Petal system, 149
- Petri nets, 22

- Ph.D. dissertation, 134
- Physical clocks, 190–192
- Physical considerations, 229–231
- PlusCal language, 18, 119–122, 124, 159
- Pnueli’s temporal logic, 105–106
- Post-quantum cryptosystems, 81
- Practical problems, 2
- Precedes relations, 103
- Preceding events, 230
- Premature termination, 253–254
- Presidential selection requirement, 293–294, 297–298
- Principles of Distributed Computing
 - Influential Paper Award, 6
- Private variables, 240
- Process fairness, 14
- Process templates, 120
- Processes
 - distributed systems, 48–49
 - mutual exclusion problem, 239–241
- Producer-consumer synchronization, 17–23, 121–122
- Progress condition, 280
- Prophecy variables, 10, 113
- Proposed decrees, 296
- Public-key cryptography, 77–78

- Quantum-safe cryptosystems, 81
- Quorums, 282

- r -bounded waiting, 253
- Random malfunctions, 225
- Rashid, Rick, 155
- Read command, 305–306
- Read executions, 232–233
- Read-only commands, 99
- Reads, 304–306
- Real-world problems, 2
- Receive actions, 312, 314–315
- Reduced Instruction Set Computers (RISC), 151–152
- Refinement concept, 104–105, 111–112
- Registers, 4–5, 35–43
- Regular registers, 35

- Reid, Brian, 143
- Relativity, 138, 229
- Release commands, 188
- Reliable systems, 69, 221–225
- Replicated state machines, introduction, 86
- Request queues, 187–188
- Request resources, 187–188
- Request service operations, 239
- RFC 677 database, 83–84
- RISC (Reduced Instruction Set Computers), 151–152
- Rodeheffer, Tom, 158–159, 167
- Rounds, 91–95, 97–98
- RSA cryptosystem, 80–81

- Safe registers, 35, 43
- Safety properties, 10, 14, 104–105, 127
- Saxe, Jim, 152
- Schneider, Fred, 86, 139, 166–167, 170
- Schroeder, Mike, 149, 155
- Scribe document production systems, 143–145
- Secure digital signatures, 79
- Self-stabilizing algorithms, 258, 263
- Self-stabilizing multidecree parliament, 303
- Send actions, 312, 315
- Sender processes, 69
- Sequential consistency, 5, 30, 44–46, 143, 197–198
- Shared memory, 5, 29–46, 147
- Shared registers, 4–5
- Shared variables, 35–36
- Shopkeeper math, 131
- Shostak, Robert, 7–8, 141
- Shutdown behavior, 254
- Shutdown safe algorithms, 254
- SIFT (Software Implemented Fault Tolerance) project, 67, 140
- Sign operation, 78–80
- Signed messages, 213–216
- SIGOPS Hall of Fame Award, 162
- Simonyi, Charles, 143
- Single-decree Synod, 281–294, 307–311
- Slow reads, 304–306

- SMR (state machine replication) algorithm, 6, 54–58
- SMTsolvers, 126
- SMV model checker, 123
- Software Implemented Fault Tolerance (SIFT) project, 67, 140
- Software specification, 135
- Space-like planes, 230
- Space-time diagrams, 181–183
- Special relativity, 229
- SPHINCS scheme, 81
- Spin model checker, 123
- SRI years (1977–1985), 139–146
- Stable property, 64
- Standard model, 15
- Starvation, 31–32
- State machine replication (SMR) algorithm, 54–58
- State machine replication (SMR) paradigm, 6
- State machine replication with benign failures, 8–9, 83–84
 - active vs. passive replication, 85
 - asynchronous consensus protocols, 90–96
 - benign system models, 87–88
 - dynamic reconfiguration, 100–102
 - Paxos protocol, 96–100
 - protocol basics, 88–90
 - review, 85–87
- State machines, 48, 304–307
- Stockmeyer, Larry, 94
- Stoppable Paxos protocol, 102
- Strong Clock Condition, 189–191
- Strong fairness, 106
- Sturgis, Howard, 162n
- Stuttering invariance, 115, 127–129
- Stuttering steps, 105
- Synchronization, 17–23, 69, 76–77, 87, 120–122, 188, 190
- Synchronous consensus protocols, 89
- System executions, 41–42, 231–233
- System steps, 257
- Systems, concurrency, 42
- Systems Research Center (SRC), 146
- T-Agreement property, 88
- T-Completion property, 88
- T-Validity property, 88
- Taylor, Bob, 146, 155
- Temporal logic of actions (TLA), 10, 105–114, 150–151
- Terminating executions, 233
- $T_{\text{E}}X$ system, 144
- $T_{\text{E}}X82$ system, 144
- Thekkath, Chandu, 149
- Thomas, Robert, 50, 53, 137–138
- Three-Bit Algorithm, 266–269
- Ticket numbers, 15–16, 32–35
- Time in distributed systems, 47–49
- Timestamps, 5, 50–52, 54–57, 61, 137–138
- TLA (temporal logic of actions), 10, 105–114, 150–151
- TLA+ language, 115–127, 160–161
- TLA+ specification, 152–153, 156–158
- TLA+ Toolbox, 127, 159–160
- TLAPS system, 116, 125–127, 159
- TLC model checker, 123–125, 153–154, 156, 159–160
- TLP tool, 151
- Total correctness property, 10
- Total delay of messages, 191
- Total order of events, 51
- Total ordering protocols, 88
- Transient malfunctions, 256
- Transition functions, 107–108
- Trying statement, 250–252, 254, 257, 268–269, 271
- Turing Award, 164
- Turing machines, 48
- Two-arrow model, 16–17, 22–23
- Two generals paradox, 141
- Two-phase commit case, 162
- Typed languages, 115–116
- Unannounced death behavior, 255
- University of Illinois array processing computer, 135

- Unobstructed exit conditions, 32
- Unpredictable delay of messages, 191
- Update command, 305–306
- van Renesse, Robbert, 345
- Vector clocks, 52–53
- Verify operation, 78, 80
- Version vectors, 53
- Vertical Paxos protocol, 101, 161
- Viewstamped replication (VR) protocol, 99–100
- Vogt, Friedrich “Fritz,” 158
- Wait-freedom condition, 4, 36–37
- Weak fairness, 106
- Weakening shared variables, 35–36
- Web 2.0, 157
- Web Services Atomic Transaction protocol, 128
- Welch, Jennifer L., 345
- Western Research Lab, 147
- Whorfian syndrome, 159n
- Wildfire multiprocessor system, 152
- World line of point objects, 229
- Writers, 40–41
- XBox 360 memory system, 128
- Xerox PARC, 139–140
- XMSS scheme, 81
- Yu, Yuan, 152–153, 166
- Zambrovski, Simon, 160
- Zenon tableau prover, 126
- Zermelo-Fraenkel set theory with choice (ZFC), 115
- Zhou, Lidong, 161, 165

Biographies

Dahlia Malkhi



Dahlia Malkhi received her Ph.D., an M.Sc. and a B.Sc. in computer science from the Hebrew University of Jerusalem. She is currently a lead researcher at Calibra. She is an ACM fellow (2011), serves on the Simons Institute Advisory Board, on the MIT Cryptocurrency Journal Advisory Board, and on the editorial boards of the Distributed Computing Journal. Her research career spans between industrial research and academia: 2014–2019, founding member and principal researcher at VMware Research; 2004–2014, principal researcher at Microsoft Research, Silicon Valley;

1999–2007, associate professor at the Hebrew University of Jerusalem; 1995–1999, senior researcher, AT&T Labs-Research, New Jersey. She has research interest in applied and foundational aspects of reliability and security in distributed systems.

Authors

Karolos Antoniadis is a Ph.D. candidate at EPFL under the supervision of Prof. Rachid Guerraoui. He holds an MSc in Computer Science from ETH Zurich.

Hagit Attiya received the B.Sc. degree in Mathematics and Computer Science from the Hebrew University of Jerusalem, in 1981, the M.Sc. and Ph.D. degrees in Computer Science from the Hebrew University of Jerusalem, in 1983 and 1987, respectively. She is a professor at the department of Computer Science at the Technion, Israel Institute of Technology, and holds the Harry W. Labov and Charlotte Ullman Labov Academic Chair. Before joining the Technion, she has been a post-doctoral

research associate at the Laboratory for Computer Science at MIT. Her research interests are in distributed and concurrent computing. She is the editor-in-chief of Springer's journal *Distributed Computing* and a fellow of the ACM.

Christian Cachin is a professor of computer science at the University of Bern, where he leads the cryptology and data security research group since 2019. Before that he worked for IBM Research—Zurich during more than 20 years. With a background in cryptography, he is interested in all aspects of security in distributed systems, particularly in cryptographic protocols, consistency, consensus, blockchains, and cloud-computing security. He contributed to several IBM products, formulated security standards, and helped to create the Hyperledger Fabric blockchain platform.

Rachid Guerraoui received his Ph.D. from University of Orsay, and M.Sc from University of Sorbonne, both in Computer Science. He is currently professor in Computer Science at EPFL, Switzerland. His research interests are in the area of distributed computing.

Idit Keidar received her BSc (summa cum laude), MSc (summa cum laude), and Ph.D. from the Hebrew University of Jerusalem in 1992, 1994, and 1998, respectively. She was a Postdoctoral Fellow at MIT's Laboratory for Computer Science. She is currently a Professor at the Technion's Viterbi Faculty of Electrical Engineering, where she holds the Lord Leonard Wolfson Academic Chair. She serves as the Head of the Technion Rothschild Scholars Program for Excellence, and also heads the EE Faculty's EMET Excellence Program. Her research interests are in fault-tolerant distributed and concurrent algorithms and systems, theory and practice. Recently, she is mostly interested in distributed storage and concurrent data structures and transactions.

Roy Levin is a retired Distinguished Engineer, Managing Director, and founder of Microsoft Research Silicon Valley (2001–2014). Previously (1984–2001) he was a senior researcher and later Director of the Digital/Compaq Systems Research Center in Palo Alto, California. Before joining Digital, Levin was a researcher at Xerox's Palo Alto Research Center (1977–1984). His research focus was distributed computing systems. Levin holds a Ph.D. in Computer Science from Carnegie-Mellon University and a B.S. in Mathematics from Yale University. He is a Fellow of the Association for Computing Machinery (ACM) and a former chair of Special Interest Group on Operating Systems (SIGOPS).

Stephan Merz received his Ph.D. and habilitation in computer science from the University of Munich, Germany. He is currently a senior researcher and the head of science at Inria Nancy–Grand Est, France. His research interests are in formal veri-

fication of distributed algorithms and systems using model checking and theorem proving.

Robbert van Renesse received his Ph.D. in computer science, and an M.Sc. and a B.Sc. in mathematics from the Vrije Universiteit Amsterdam. He is currently a Research Professor at the Department of Computer Science at Cornell University and is Associate Editor of ACM Computing Surveys. He is an ACM fellow (since 2009). He was a researcher at AT&T Bell Labs in Murray Hill (1990) and served as Chair of ACM SIGOPS (2015–2019). His research interests include reliable distributed systems and operating systems.

Jennifer L. Welch received her S.M. and Ph.D. from the Massachusetts Institute of Technology and her B.A. from the University of Texas at Austin. She is currently holder of the Chevron II Professorship and Regents Professorship in the Department of Computer Science and Engineering at Texas A&M University, and is an ACM Distinguished Member. Her research interests are in the theory of distributed computing, algorithm analysis, distributed systems, mobile ad hoc networks, and distributed data structures.